

GROUP 2

- **Rozel Galceran**
- **Joeniño Cainday**
- **Angela Rose Mendez**
- **Nomeben Frietz Clarin**
- **John Aaron Sabio**



MEMORY MANAGEMENT

OBJECTIVES

01

Basic concepts of page algorithms in memory management.

02

Algorithms such as FIFO, LRU, and Page Replacement

03

Analyze the performance with different memory access patterns.

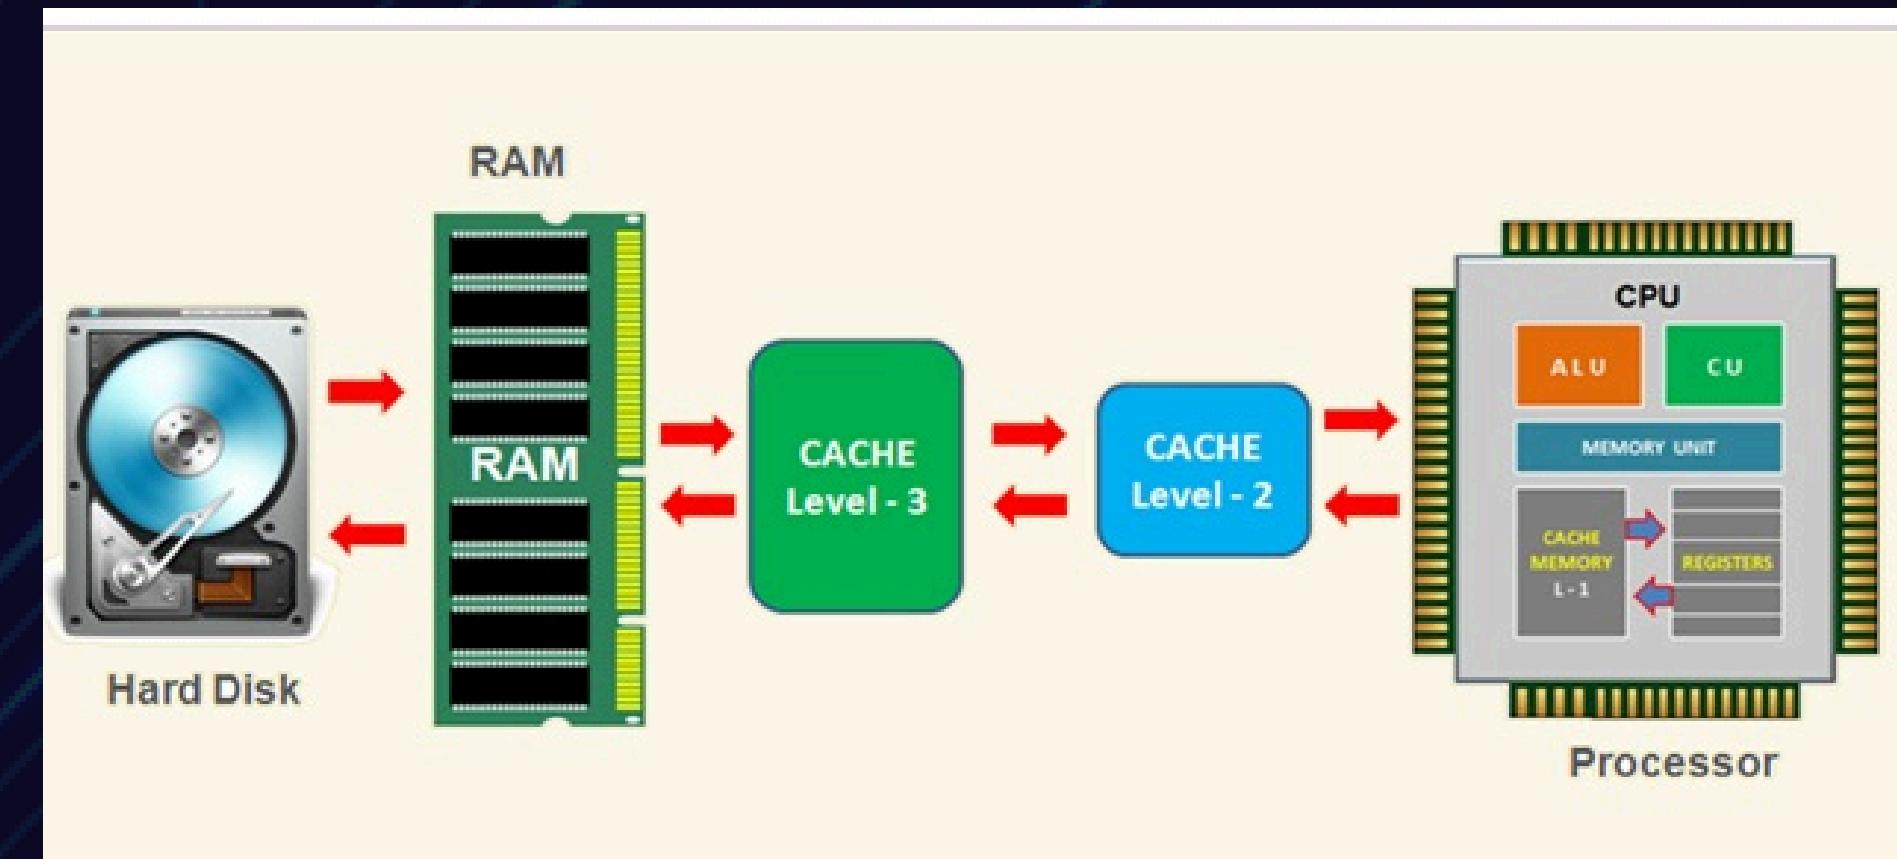




CONCEPTS

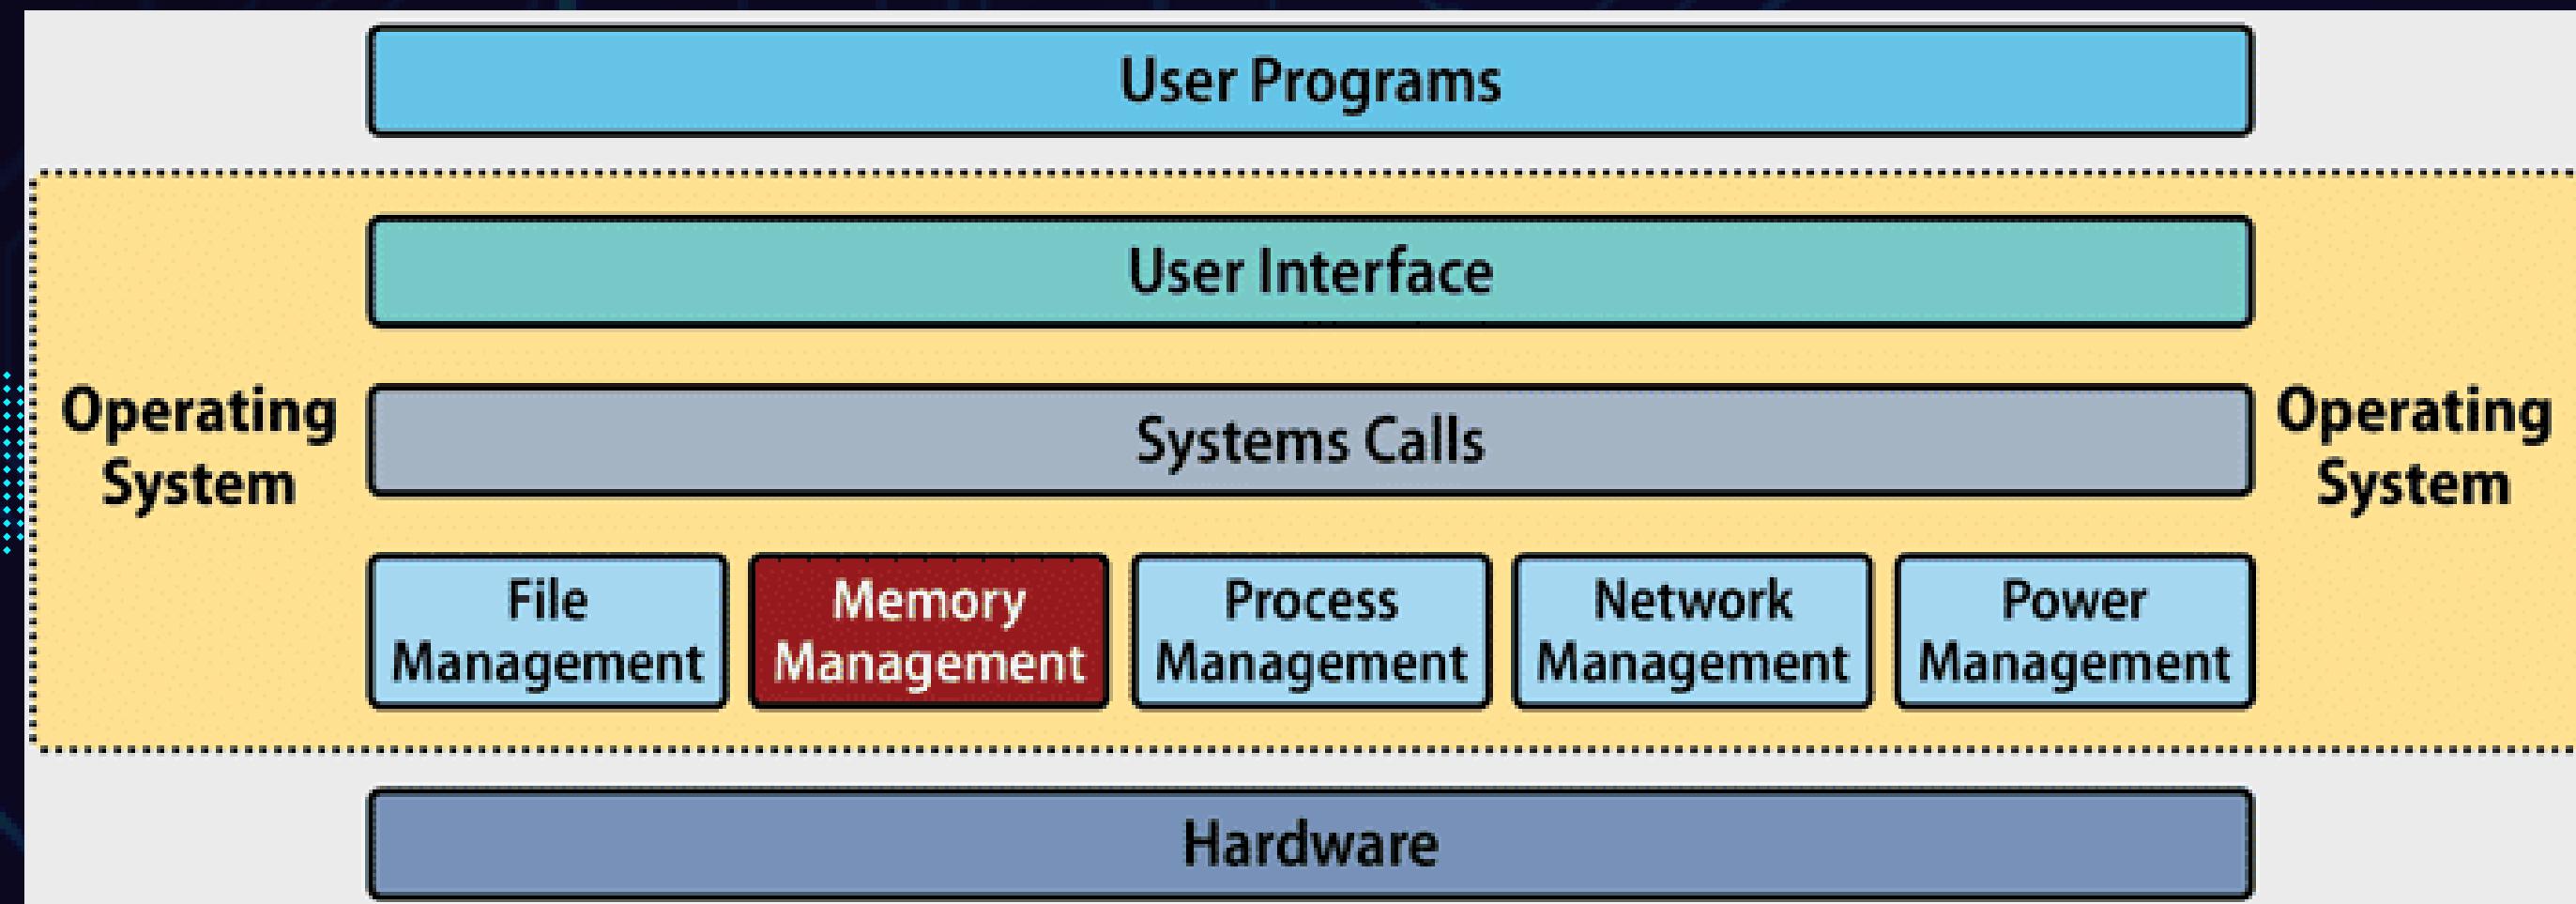
Memory Management

- a task of subdividing the memory among different processes.
- a function in the OS to manage operations between main memory and disk during process execution.
- to achieve efficient utilization of memory.



Memory Management

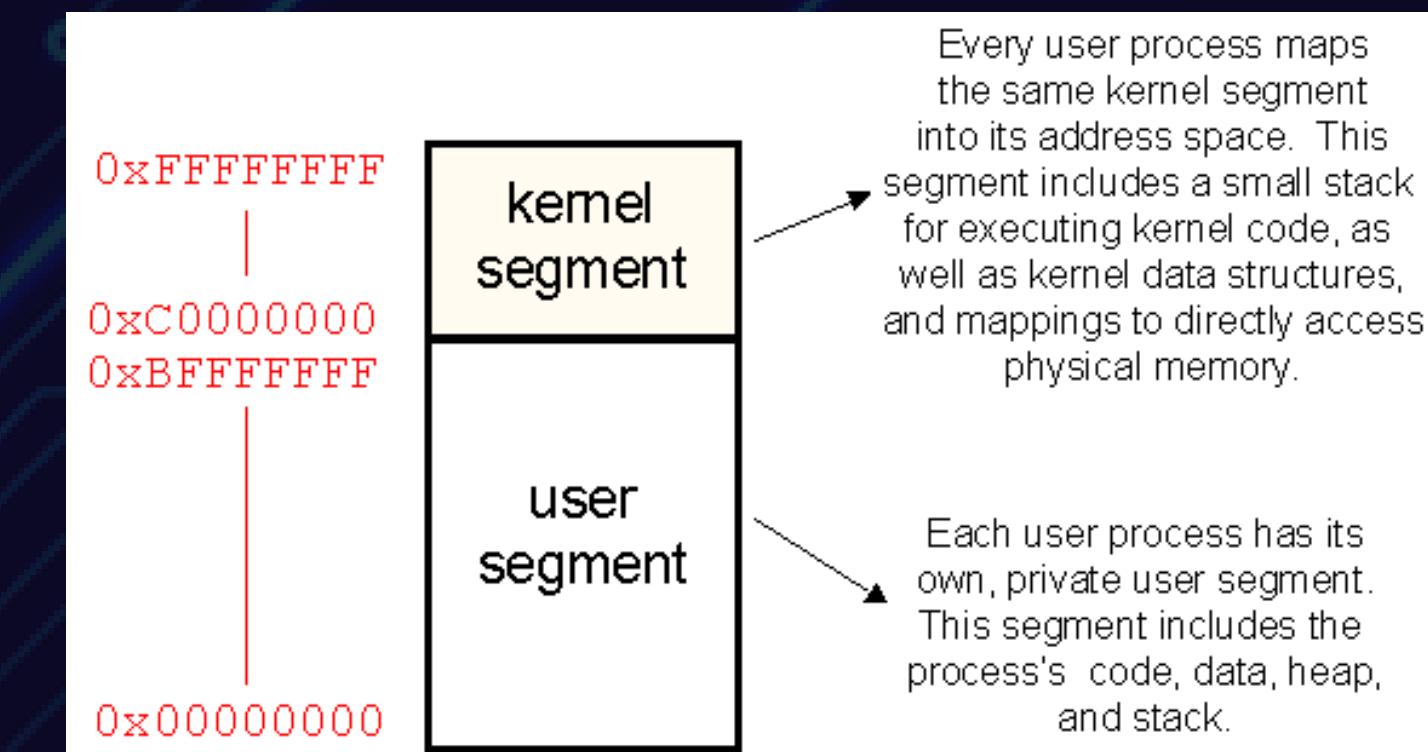
- Keeps track of each and every memory location, regardless of either it is allocated to some process or it is free.
- Checks how much memory is to be allocated to processes.
- Decides which process will get memory at what time.
- It tracks whenever some memory gets freed or unallocated and correspondingly it updates the status.



The operating system takes care of mapping the logical addresses to physical addresses at the time of memory allocation to the program.

Process Address Space

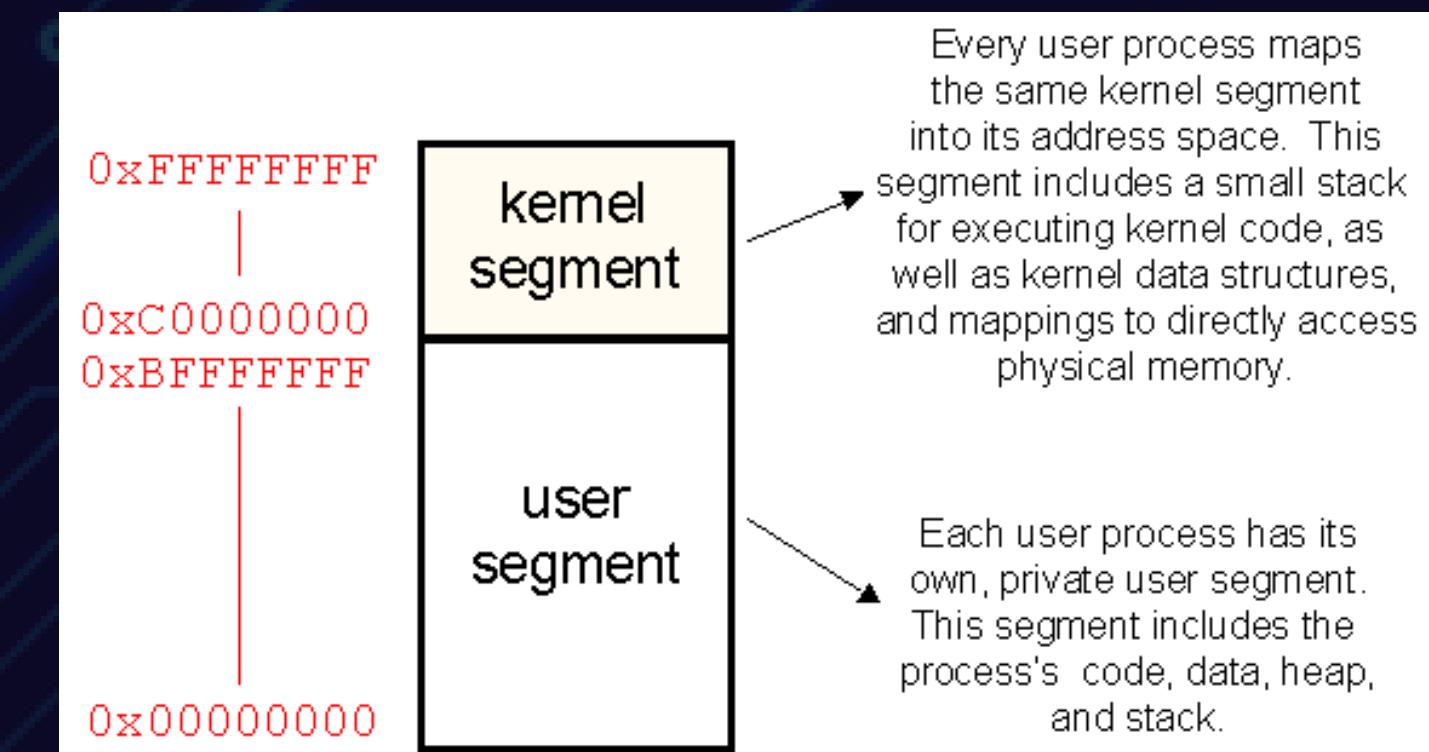
- a set of logical addresses (hex) that a process references in its code.



Process Address Space

Three types of addresses before and after memory is allocated

- Symbolic Address
 - source code variables
- Relative Address
 - compiler converted symbolic
- Physical Address
 - loader to main memory

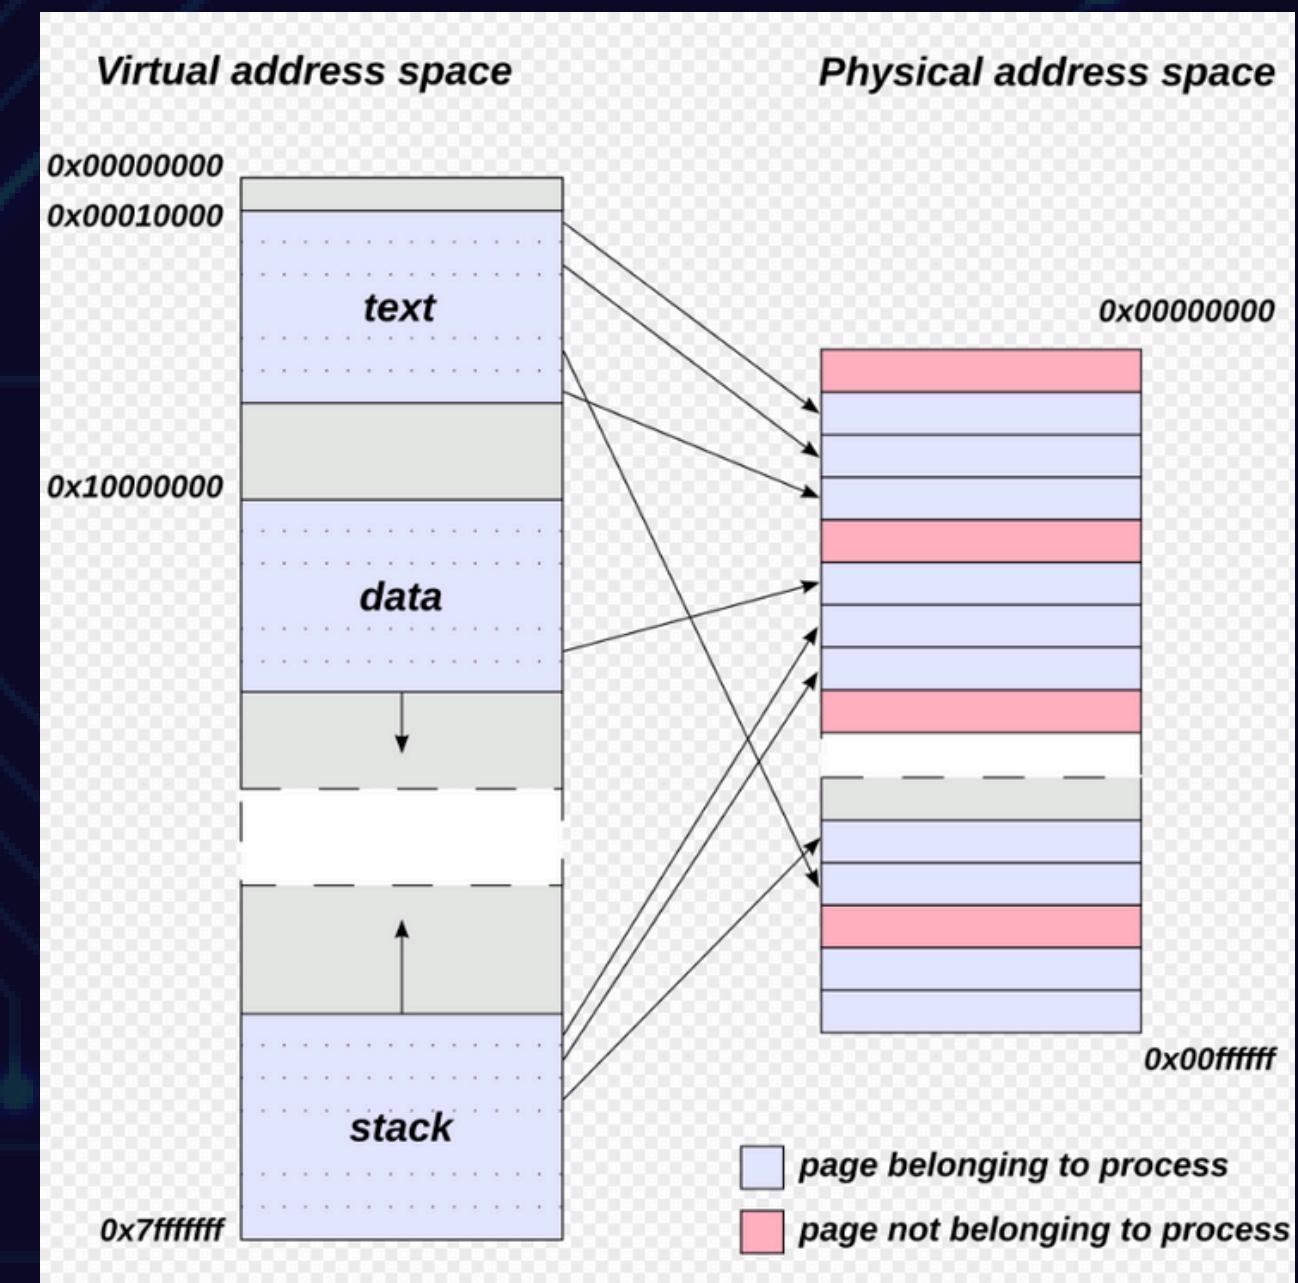


Every user process maps the same kernel segment into its address space. This segment includes a small stack for executing kernel code, as well as kernel data structures, and mappings to directly access physical memory.

Each user process has its own, private user segment. This segment includes the process's code, data, heap, and stack.

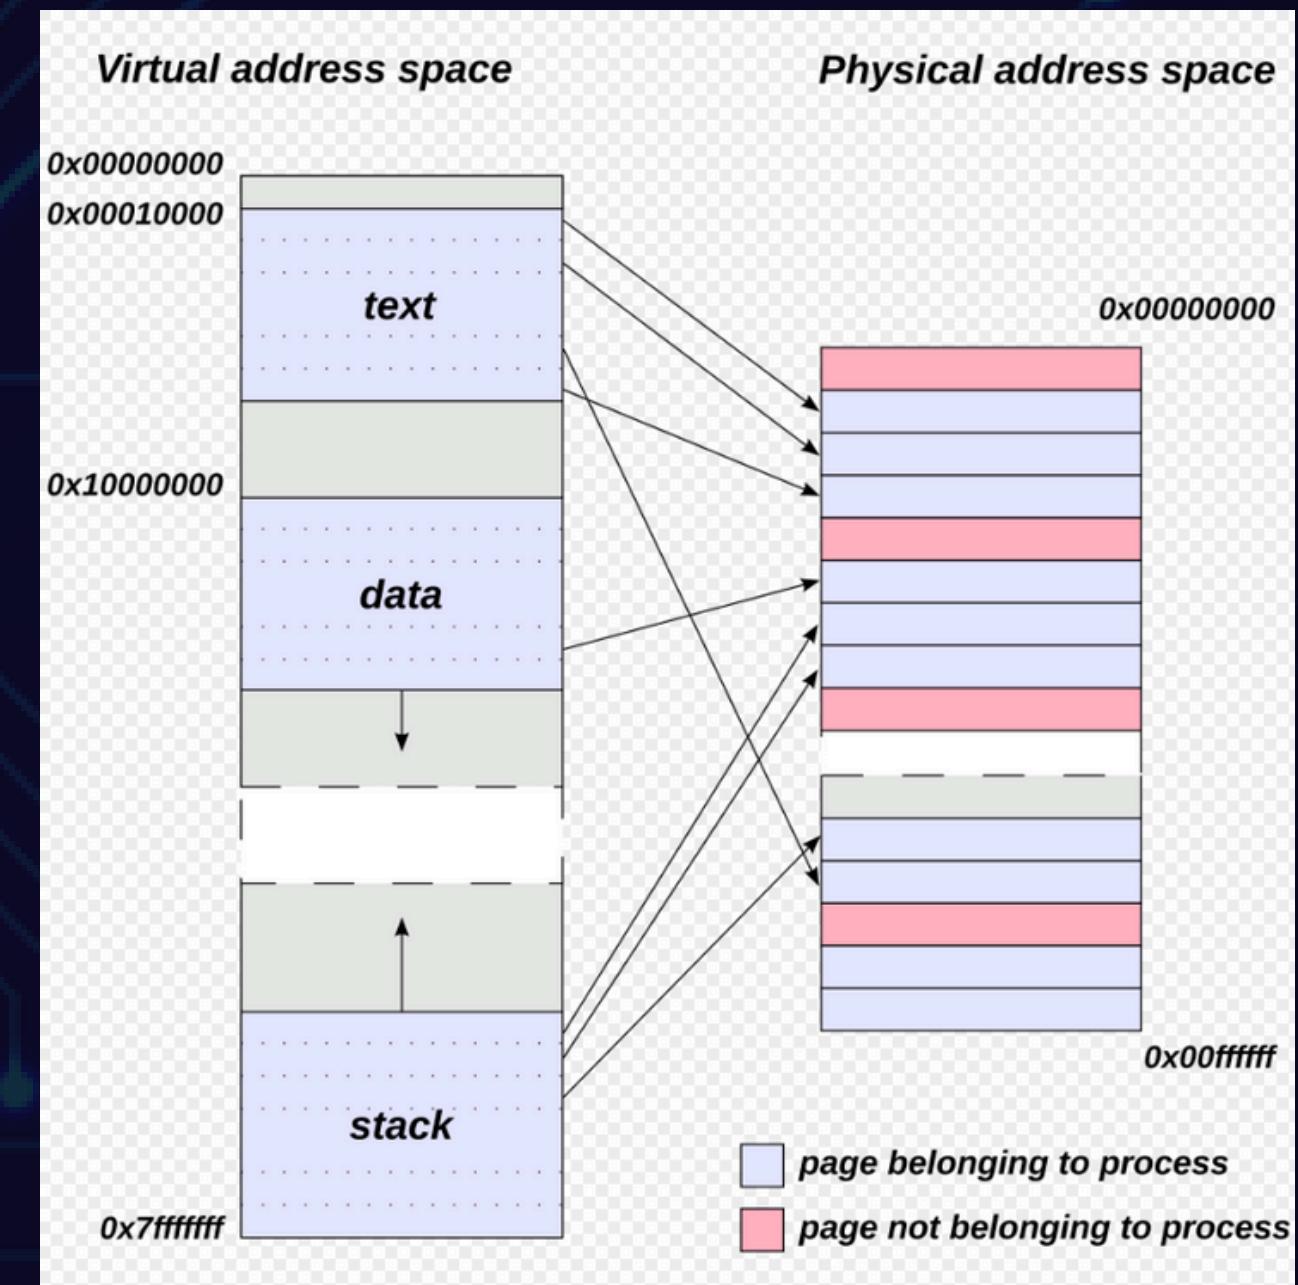
Virtual Address Space (VAS)

- the operating system set ranges of virtual addresses that makes available to a process.
 - provides isolation and prevents processes from interfering with each other



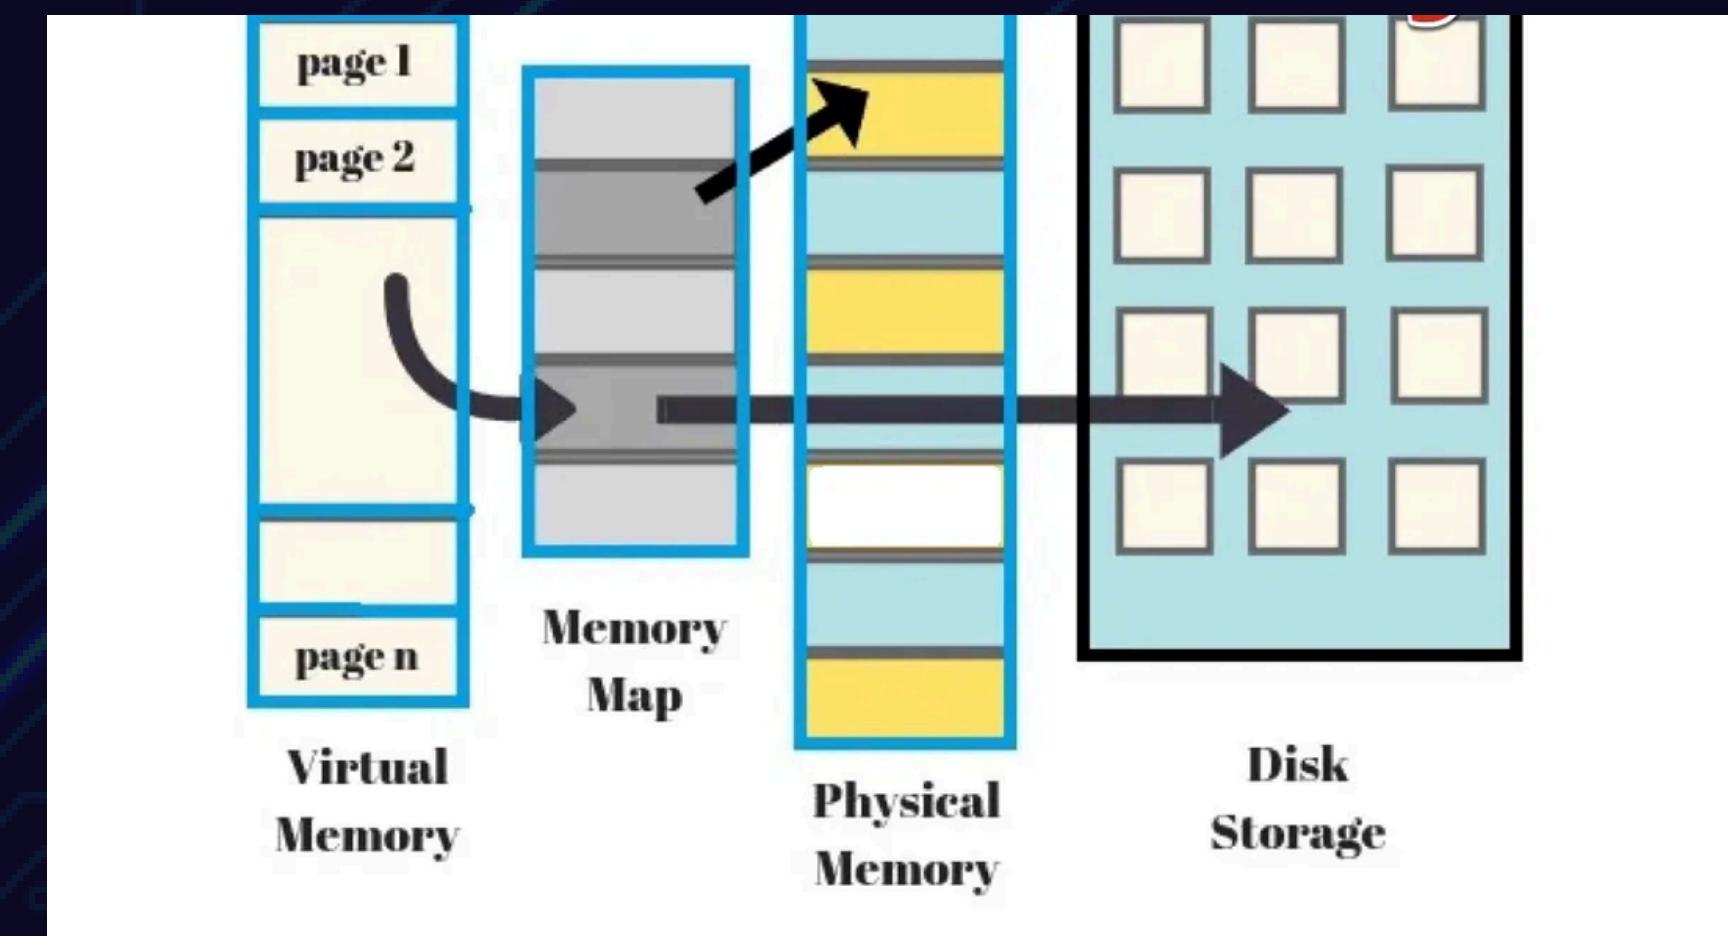
Virtual Address Space (VAS)

- aka logical address if the address generated by the CPU during the execution of a program.
 - aka virtual address when the system uses virtual memory.



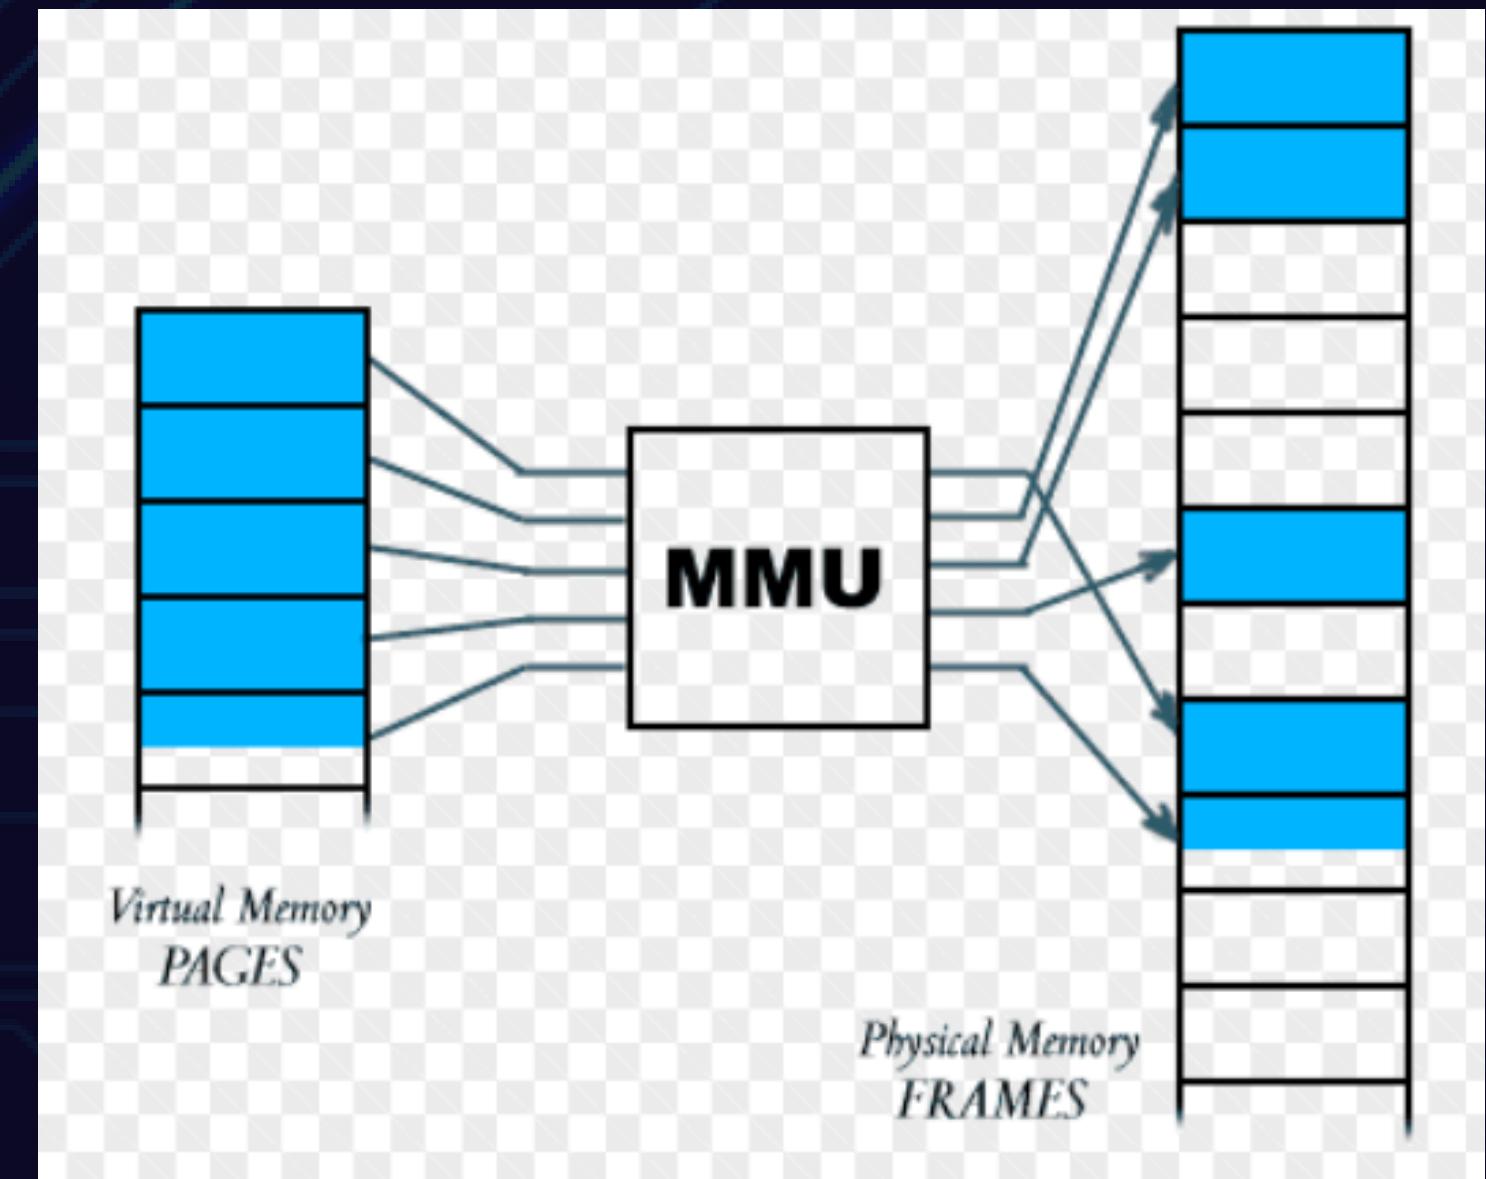
Virtual Memory

- Technique that allows the system to use disk space (secondary storage) as an extension of RAM (primary memory).
 - This allows programs to run as if there is more RAM available than physically exists on the system.



Virtual Memory

- These addresses are used by a program to access data and instructions during execution.
- However, they do not directly correspond to physical memory locations.
- Logical address is translated into a physical address by the Memory Management Unit (MMU)
- it provides an abstraction, allowing programs to think they have access to a large, continuous block of memory when, in reality, physical memory may be fragmented.

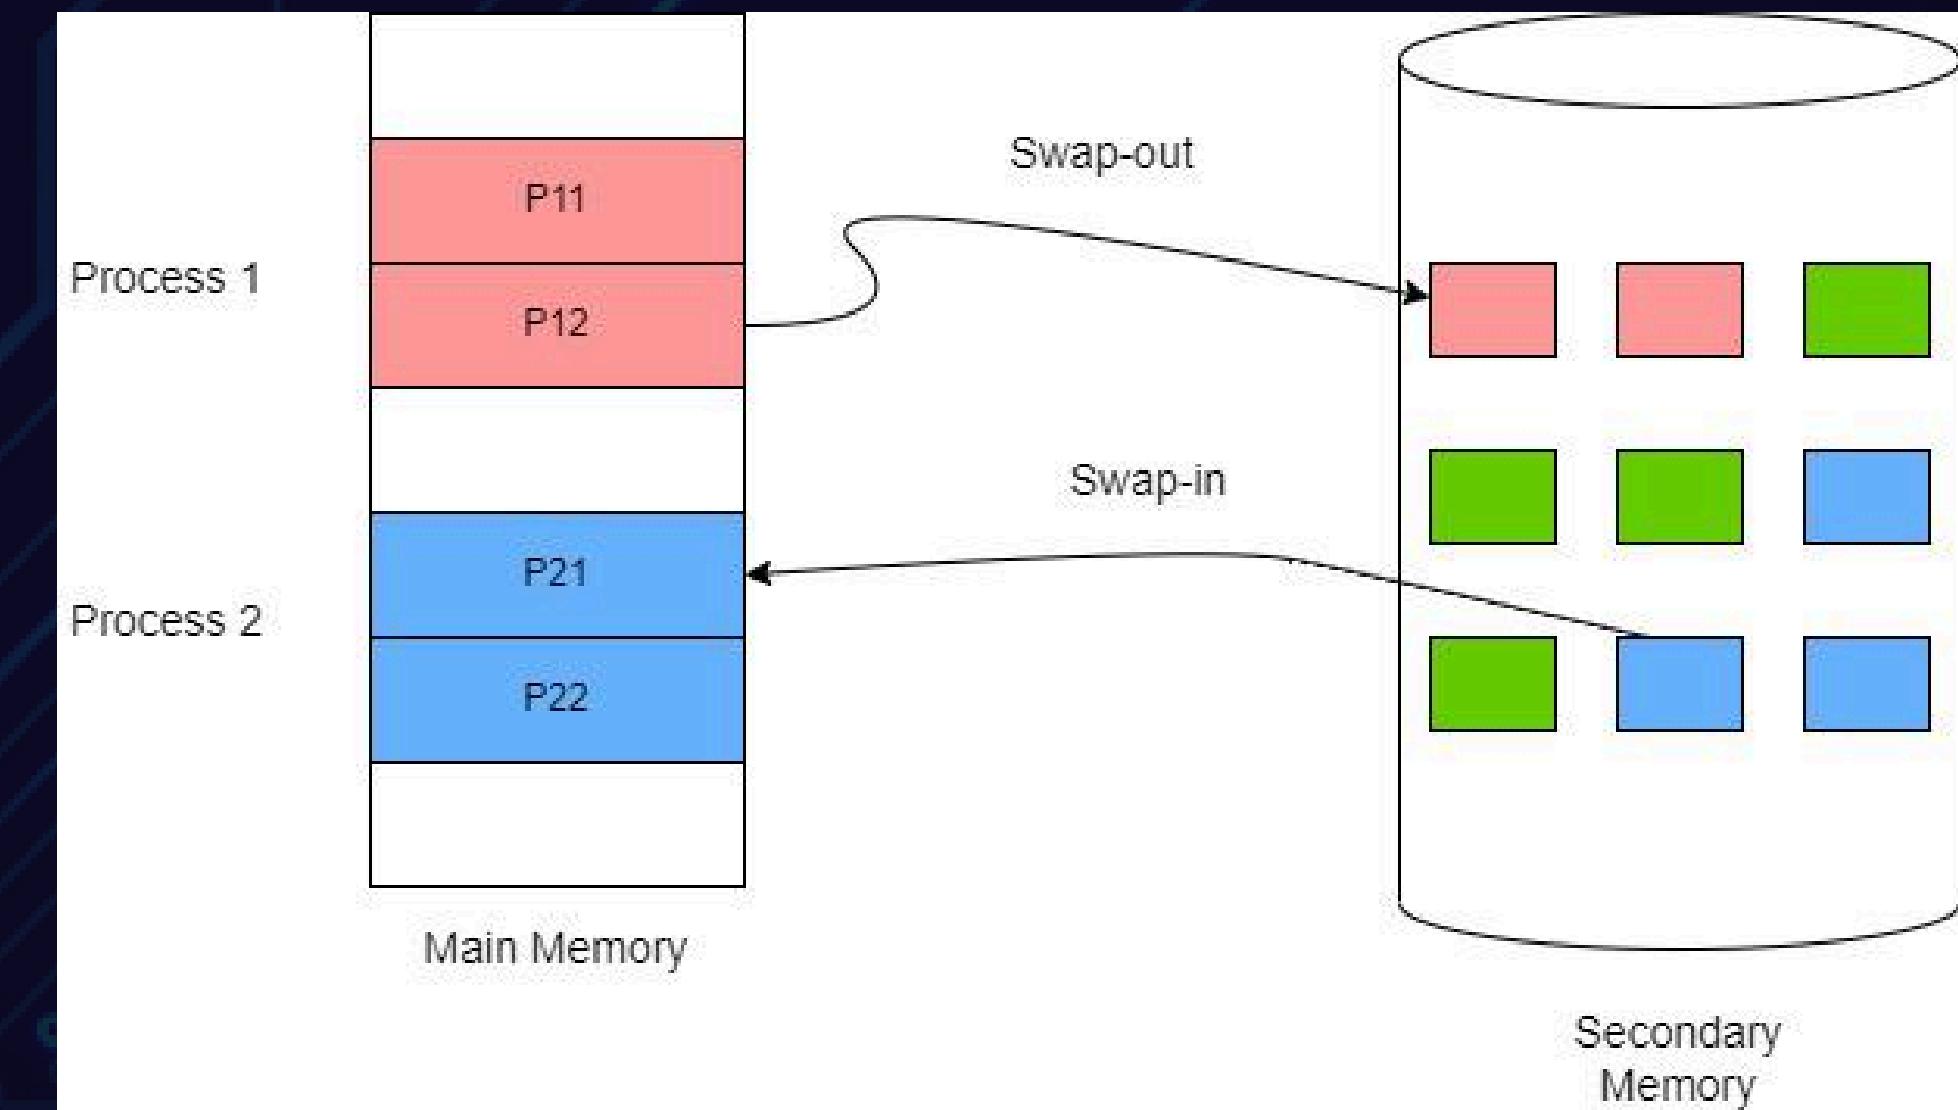




TECHNIQUES

Swapping

- a mechanism in which a process can be swapped temporarily out of main memory to disk.
 - performance is usually affected by swapping process but it helps in running multiple and big processes in parallel.

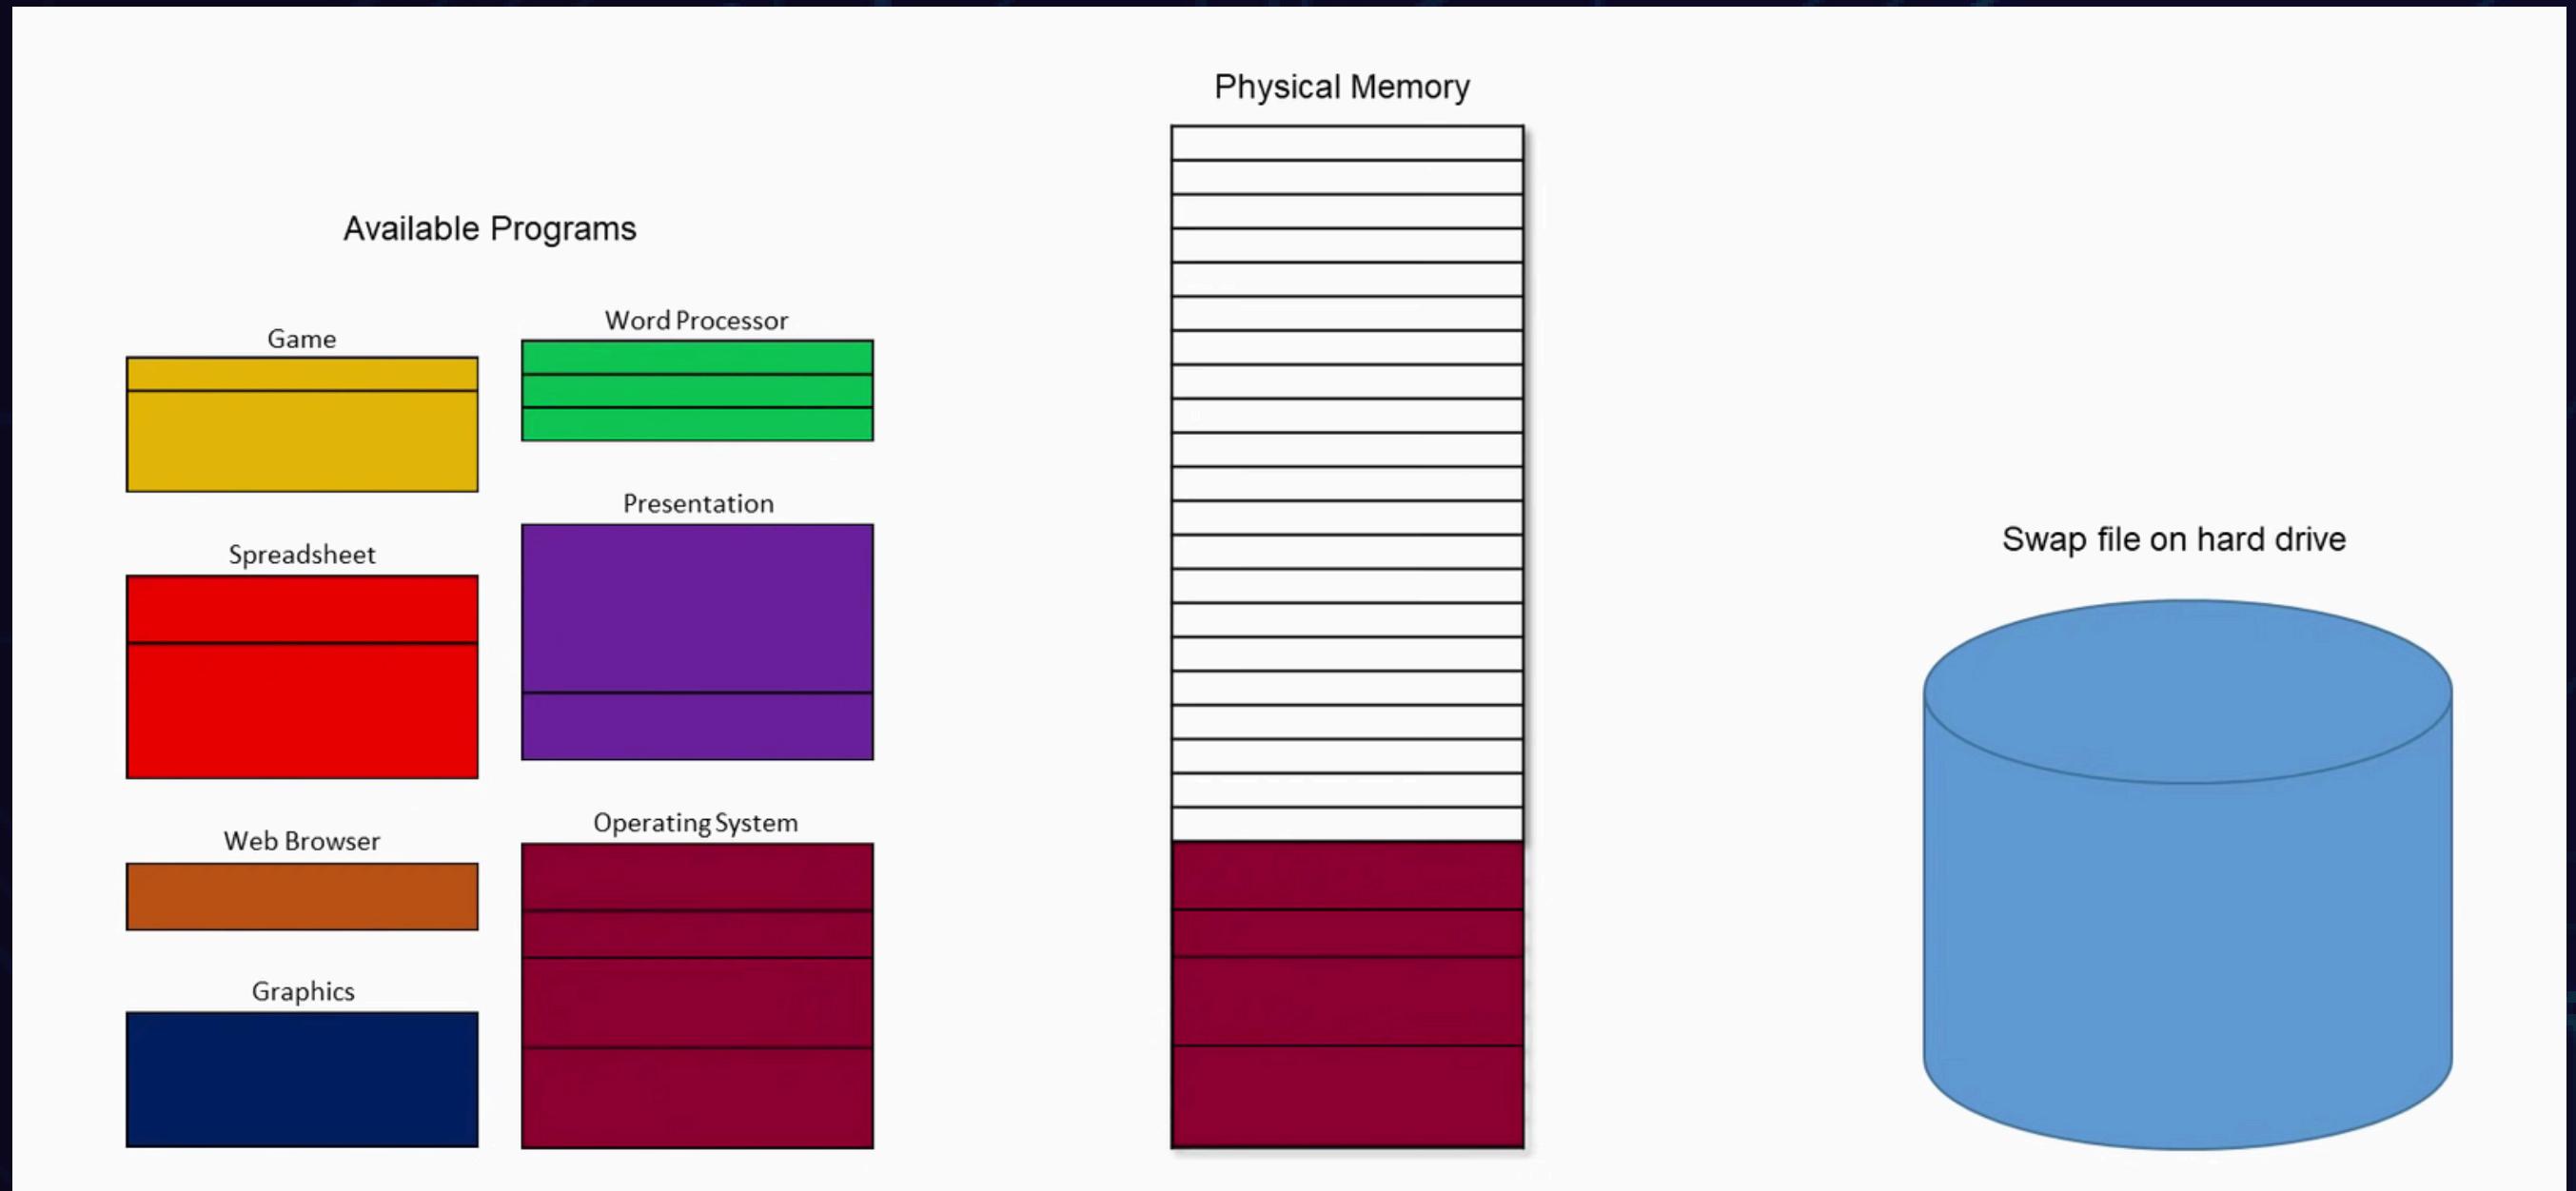


Segmentation

- A process is divided into Segments
- The chunks that a program is divided into which are not necessarily all of the exact sizes are called segments.
- Segmentation gives the user's view of the process which paging does not provide.
- The user's view is mapped to physical memory.

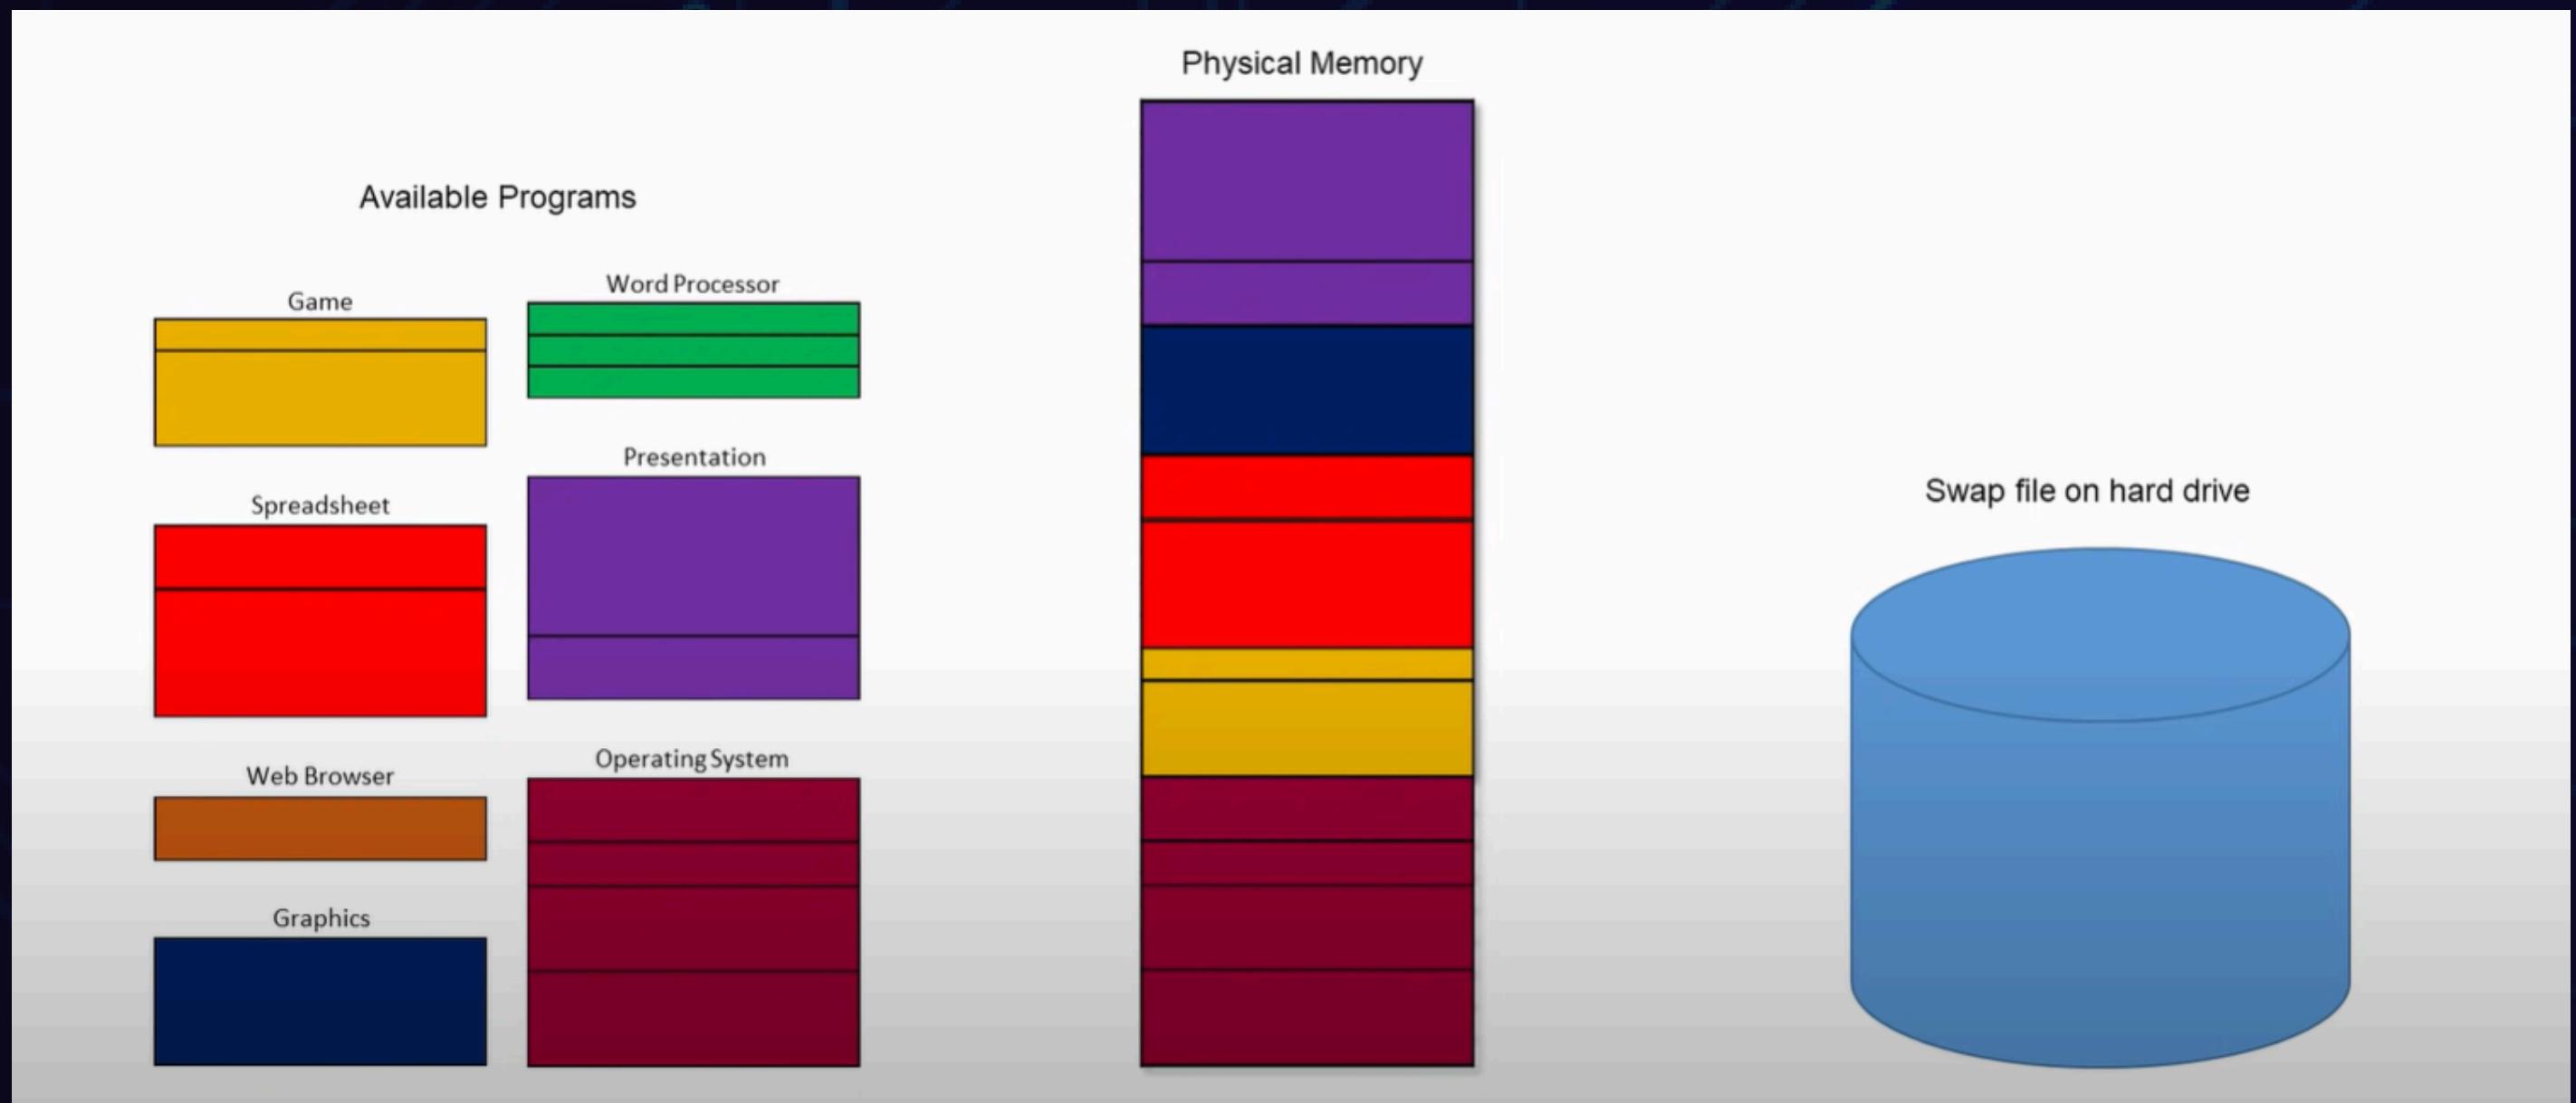


Segmentation





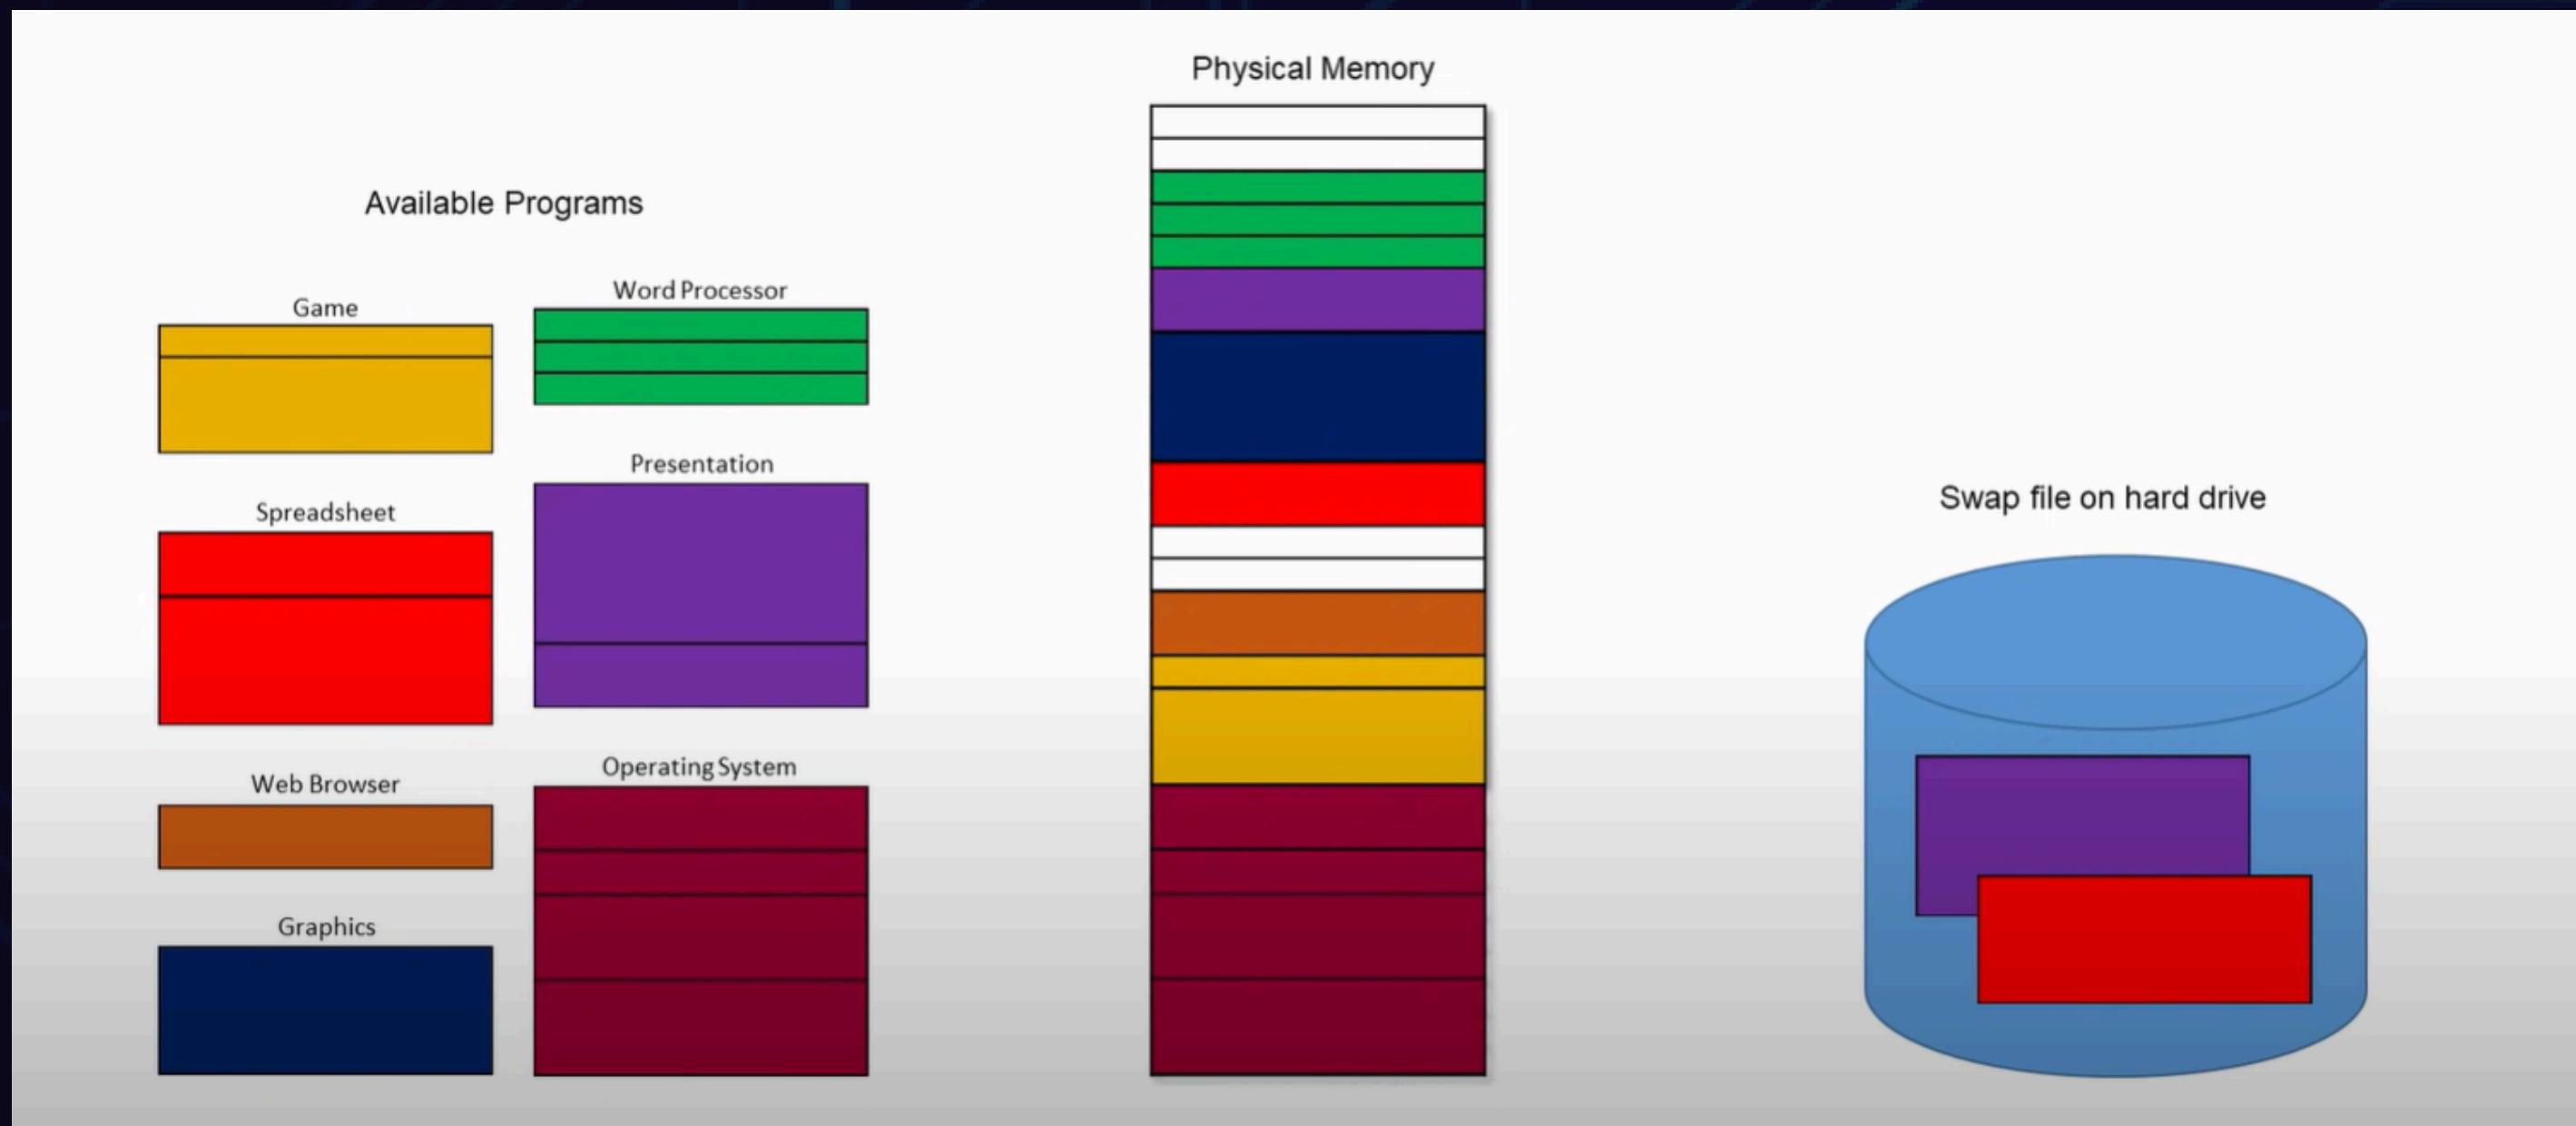
Segmentation



Types of Segmentation

- **Virtual Memory Segmentation:** Each process is divided into a number of segments, but the segmentation is not done all at once. This segmentation may or may not take place at the run time of the program.
- **Simple Segmentation:** Each process is divided into a number of segments, all of which are loaded into memory at run time, though not necessarily contiguously.

Fragmentation



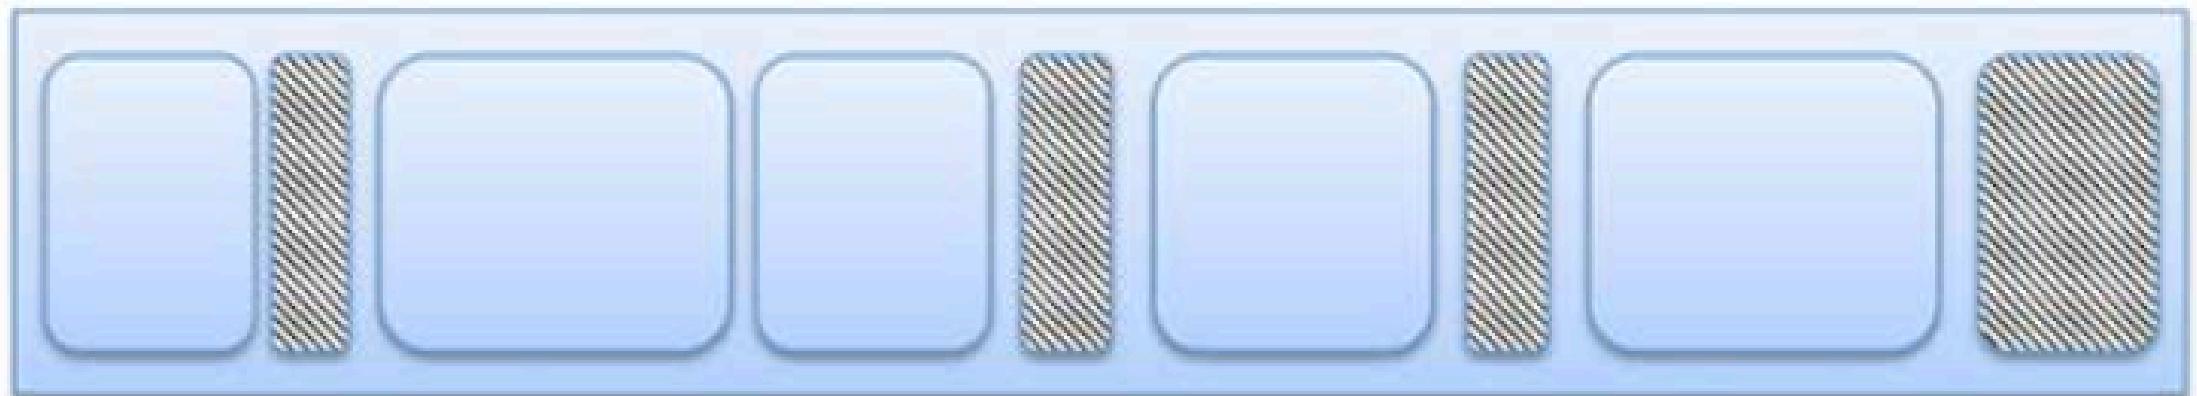
Fragmentation

- The process of dividing a computer file, such as a data file or an executable program file, into fragments that are stored in different parts of a computer's storage medium, such as its hard disc or RAM, is known as fragmentation in computing.
- When a file is fragmented, it is stored on the storage medium in non-contiguous blocks, which means that the blocks are not stored next to each other.

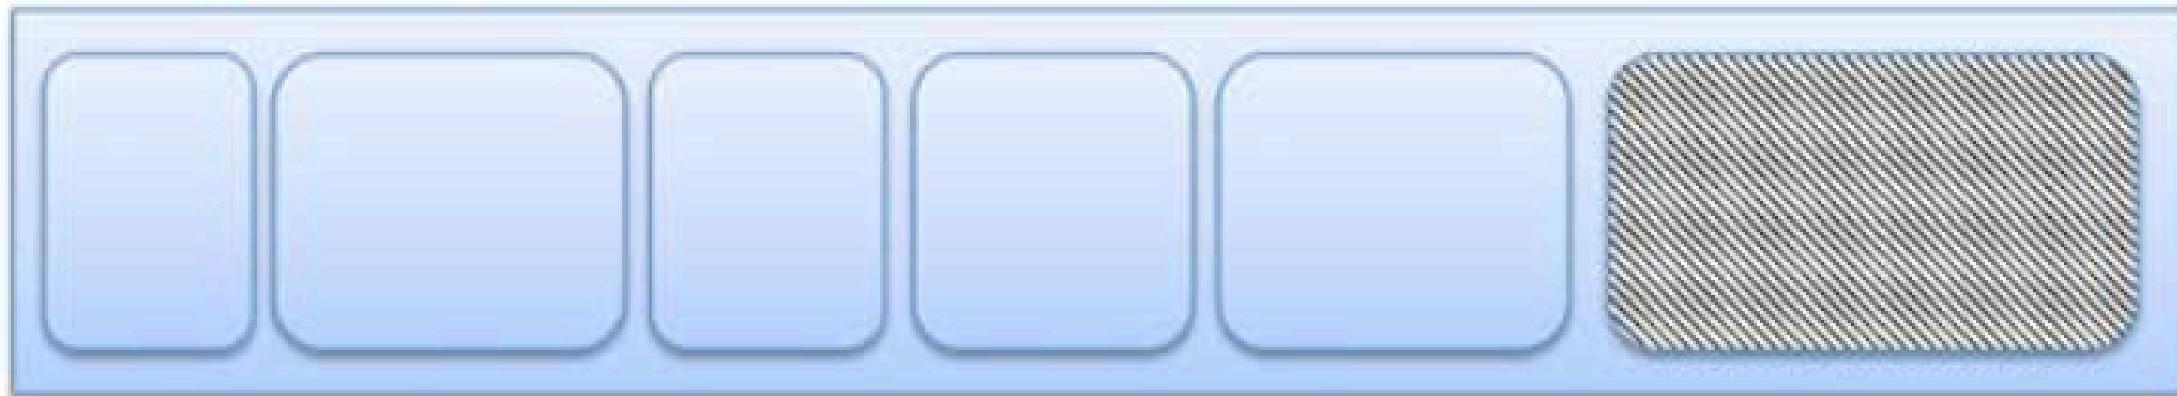


Compaction (after fragmentation)

Fragmented memory before compaction



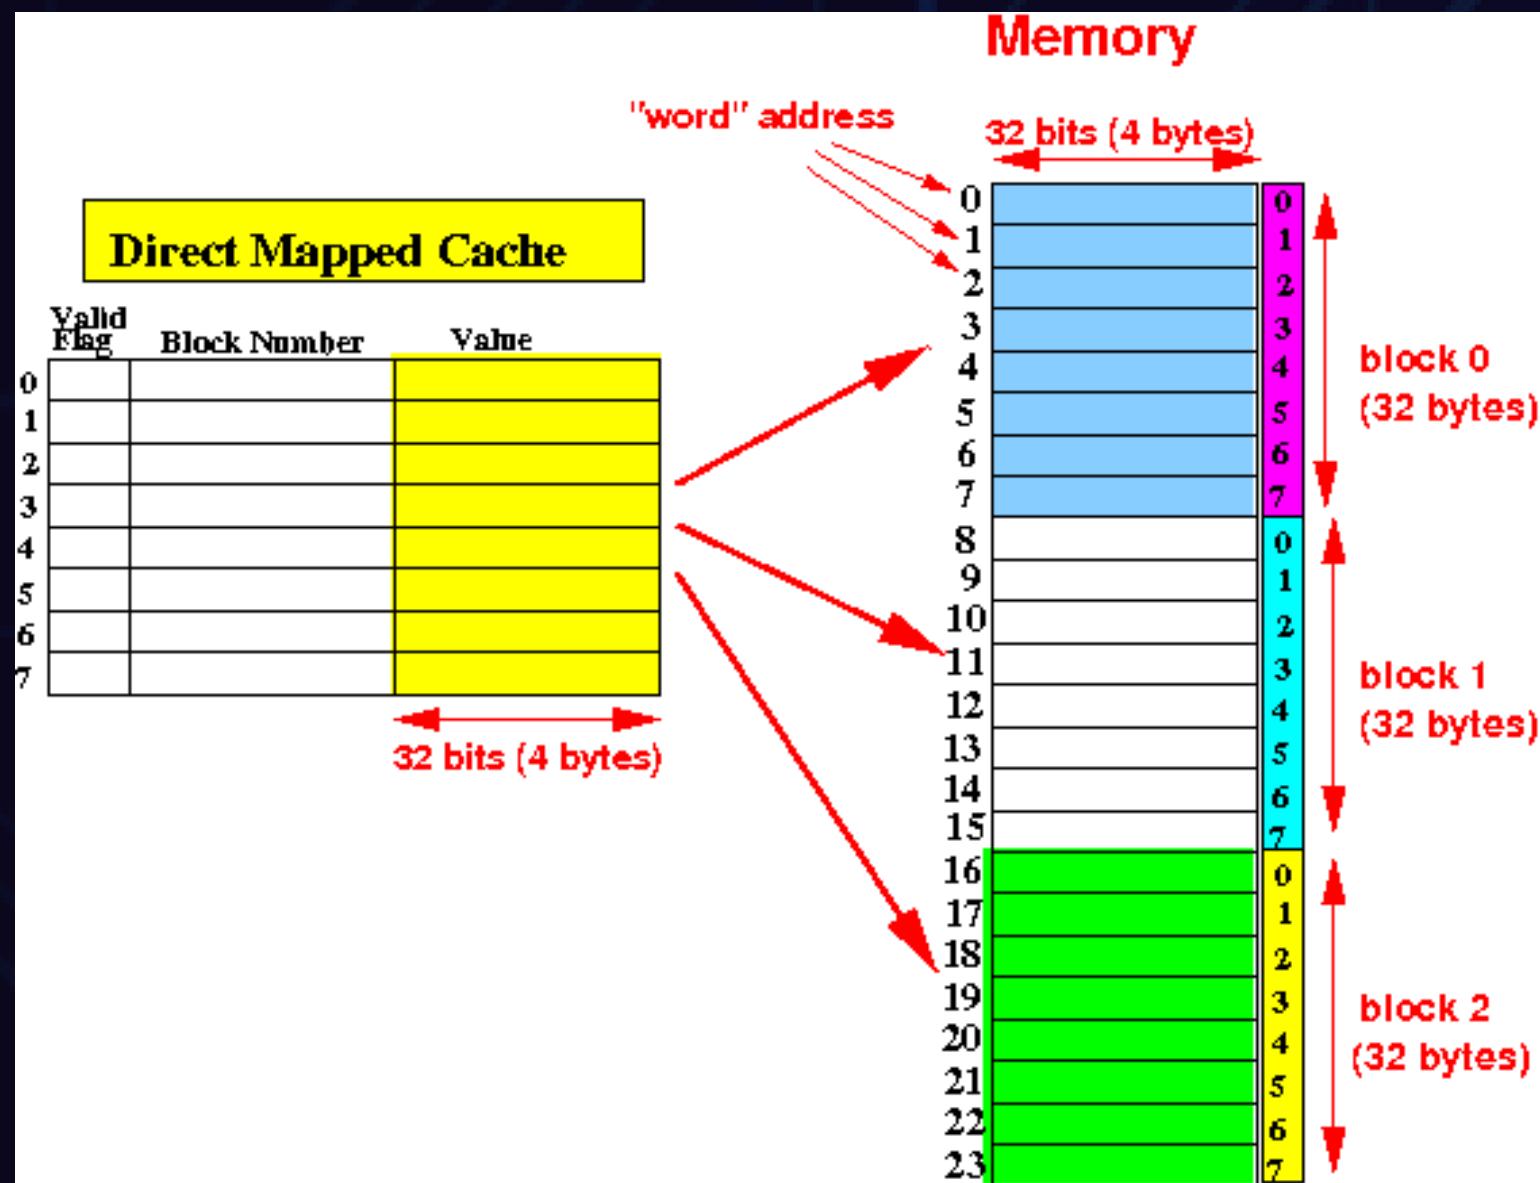
Memory after compaction



Compaction (after fragmentation)

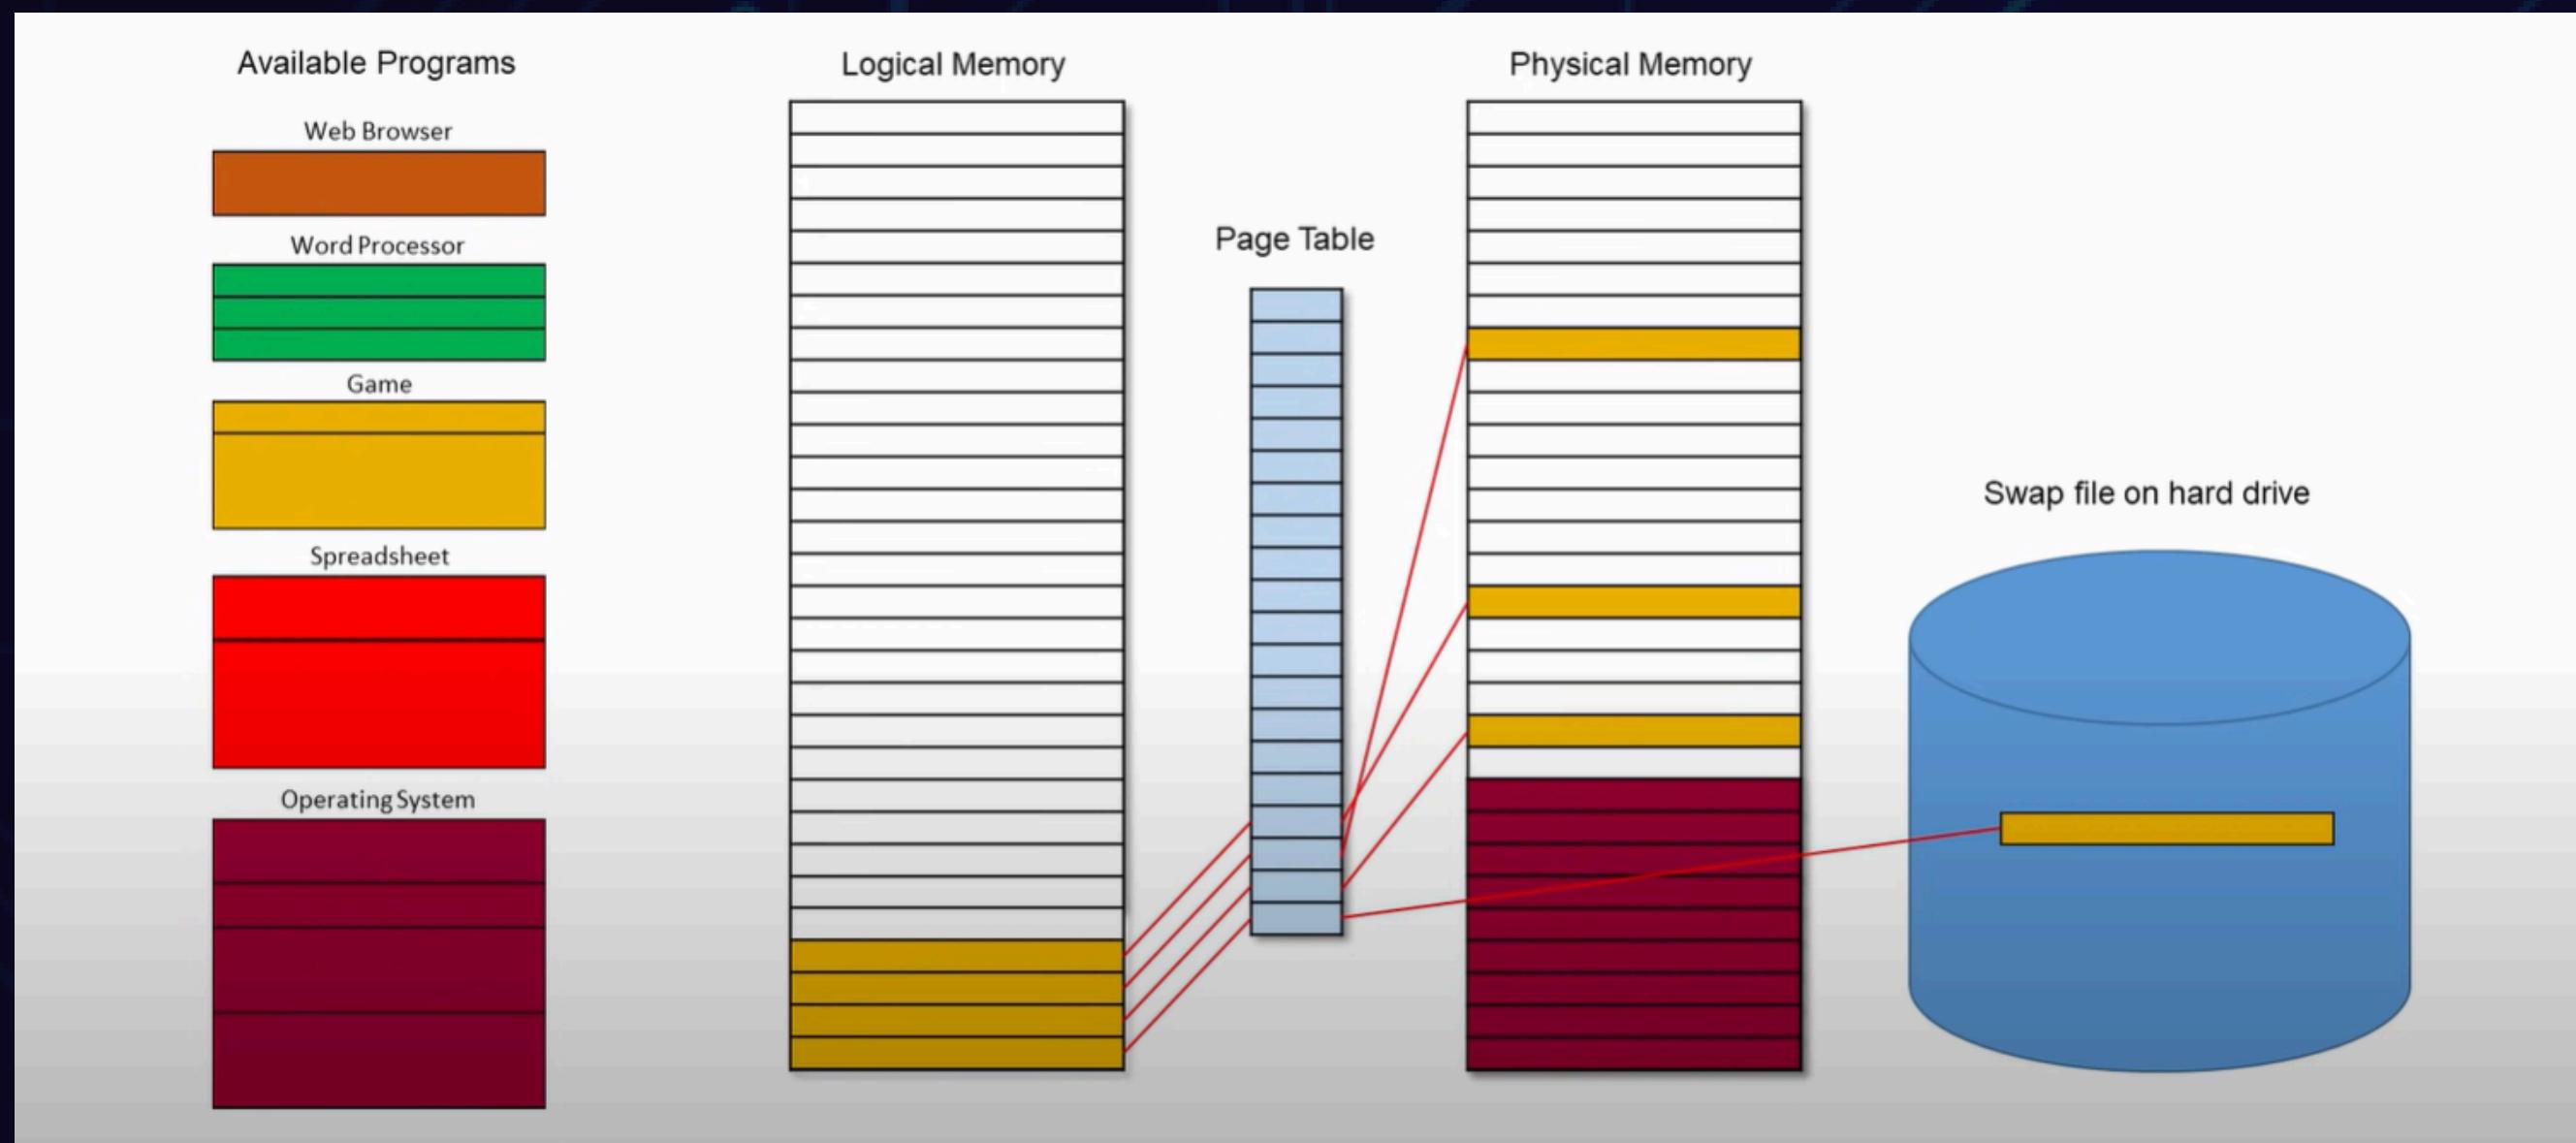
- a technique to collect all the free memory present in the form of fragments into one large chunk of free memory, which can be used to run other processes.
- It does that by moving all the processes towards one end of the memory and all the available free space towards the other end of the memory so that it becomes contiguous.

(Review Cache Lines and Blocks)



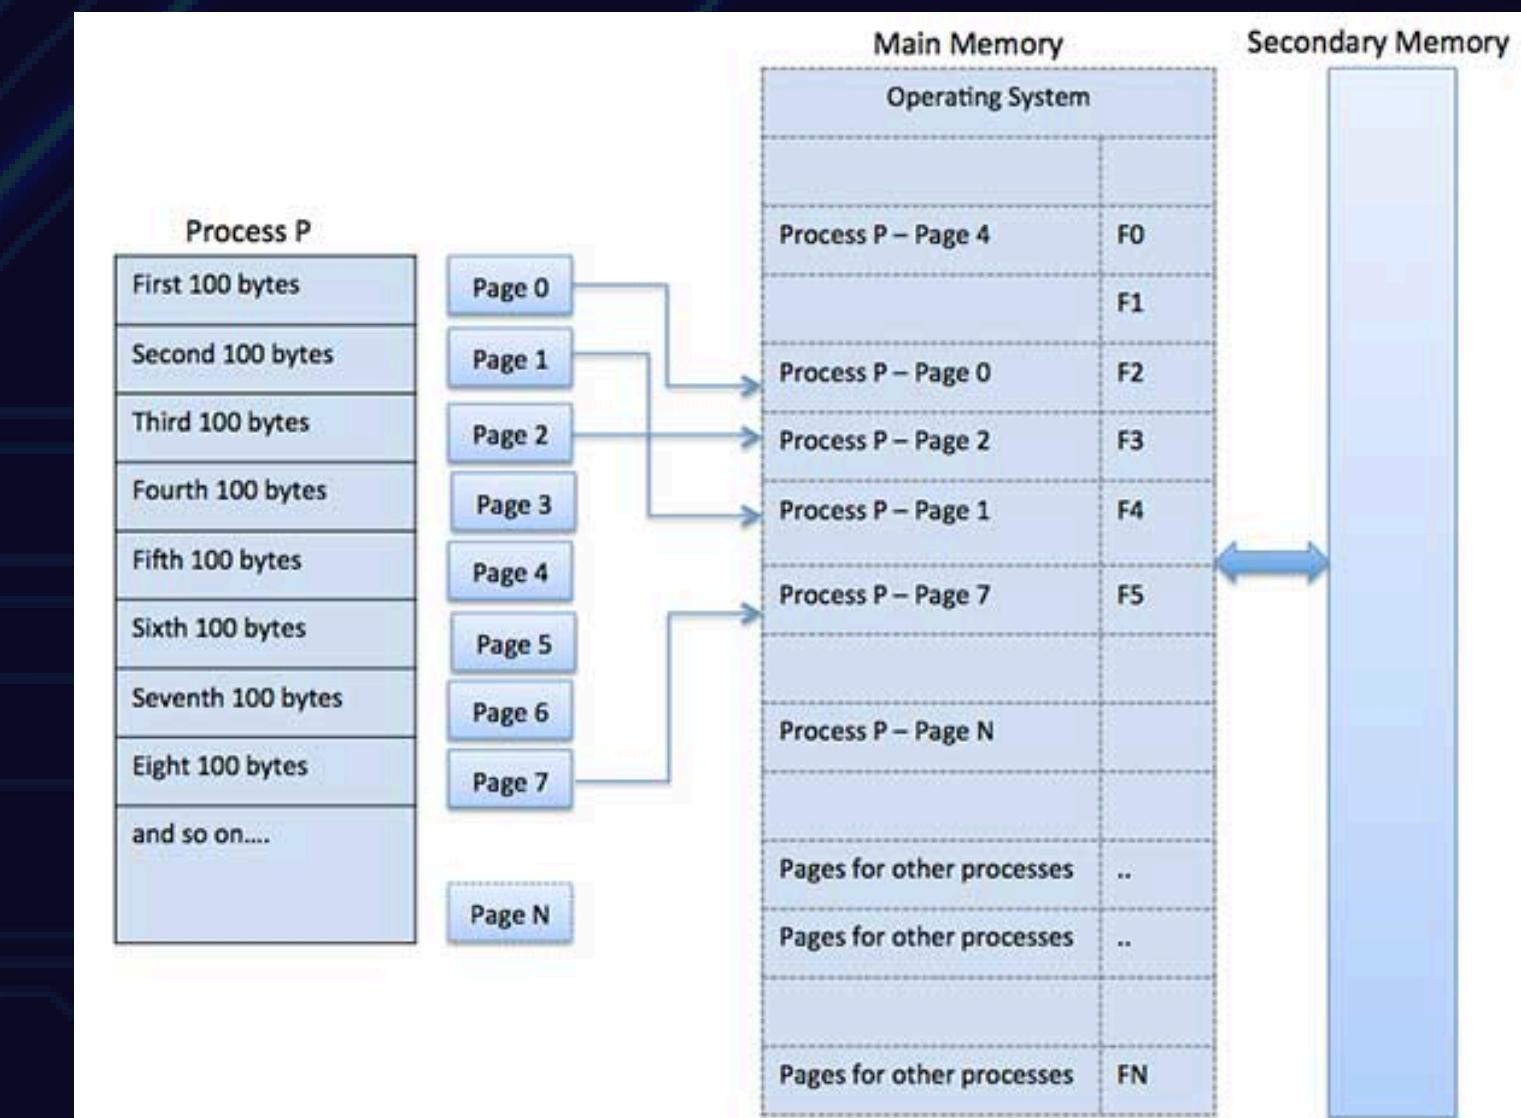
- Instead of using pages, caches operate by transferring data in smaller units called cache lines between the cache and the main memory.

Paging



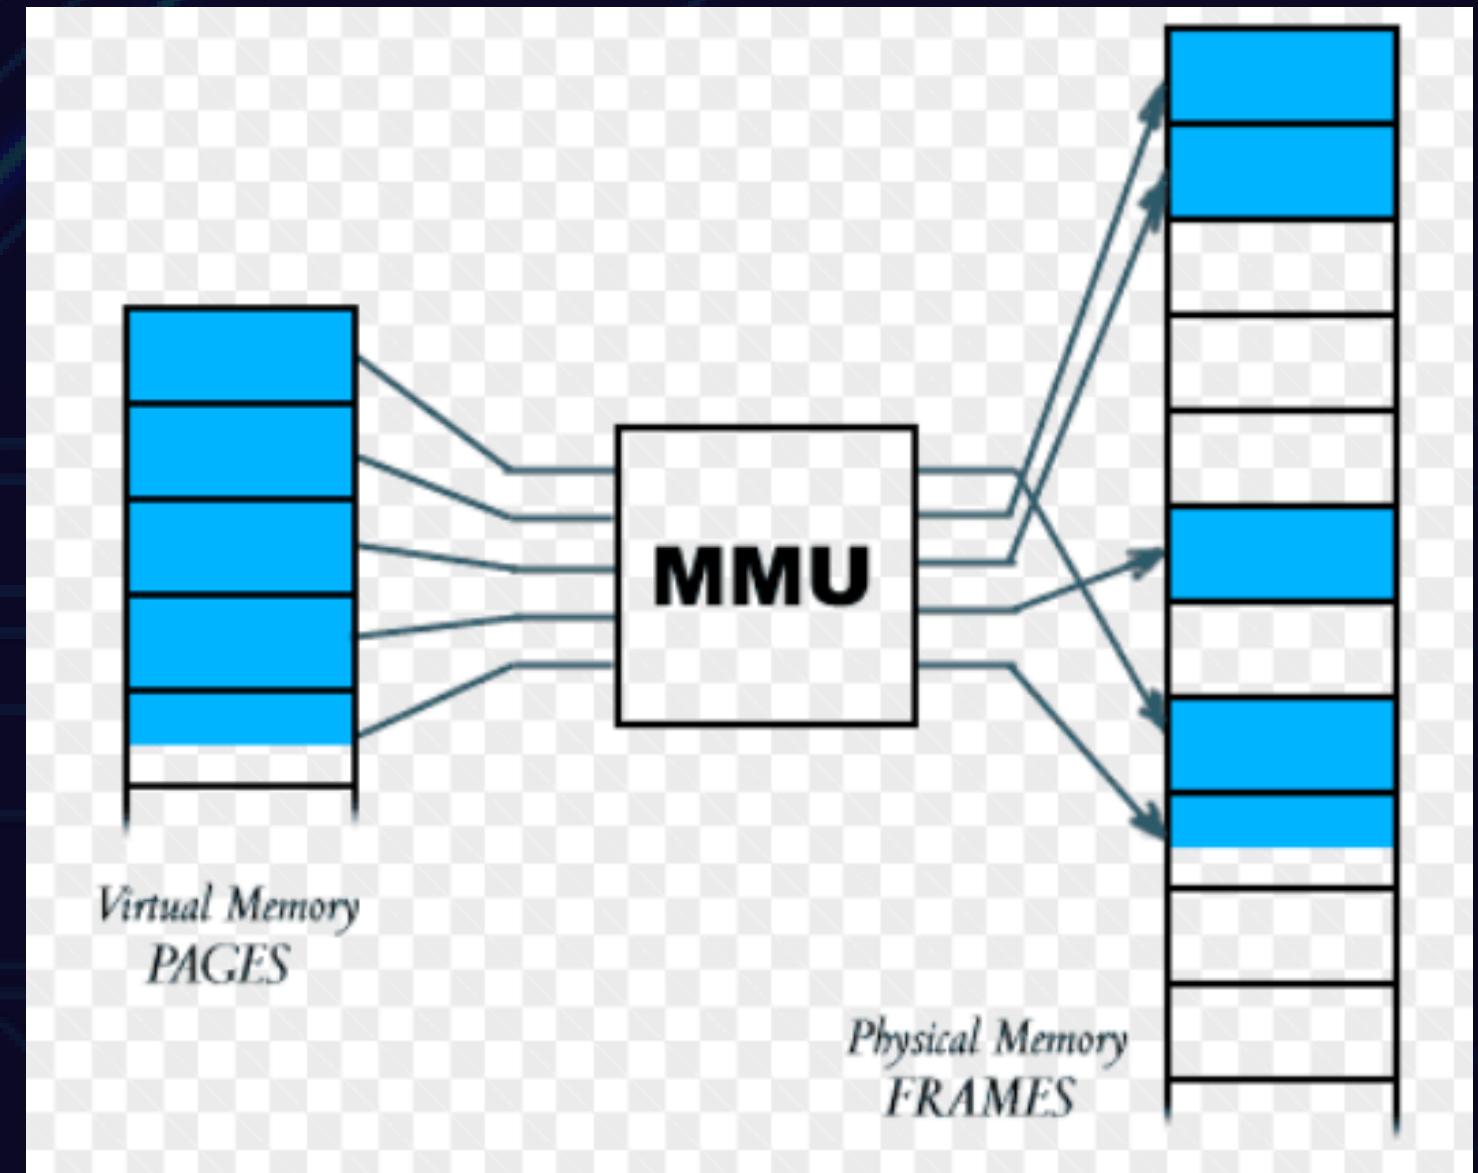
Paging

- solves the problem of segmentation and fragmentation
- **eliminates contiguous allocation** of physical memory.
- retrieving processes the form of pages from the **disk into the memory**
- to **separate each procedure** into pages.



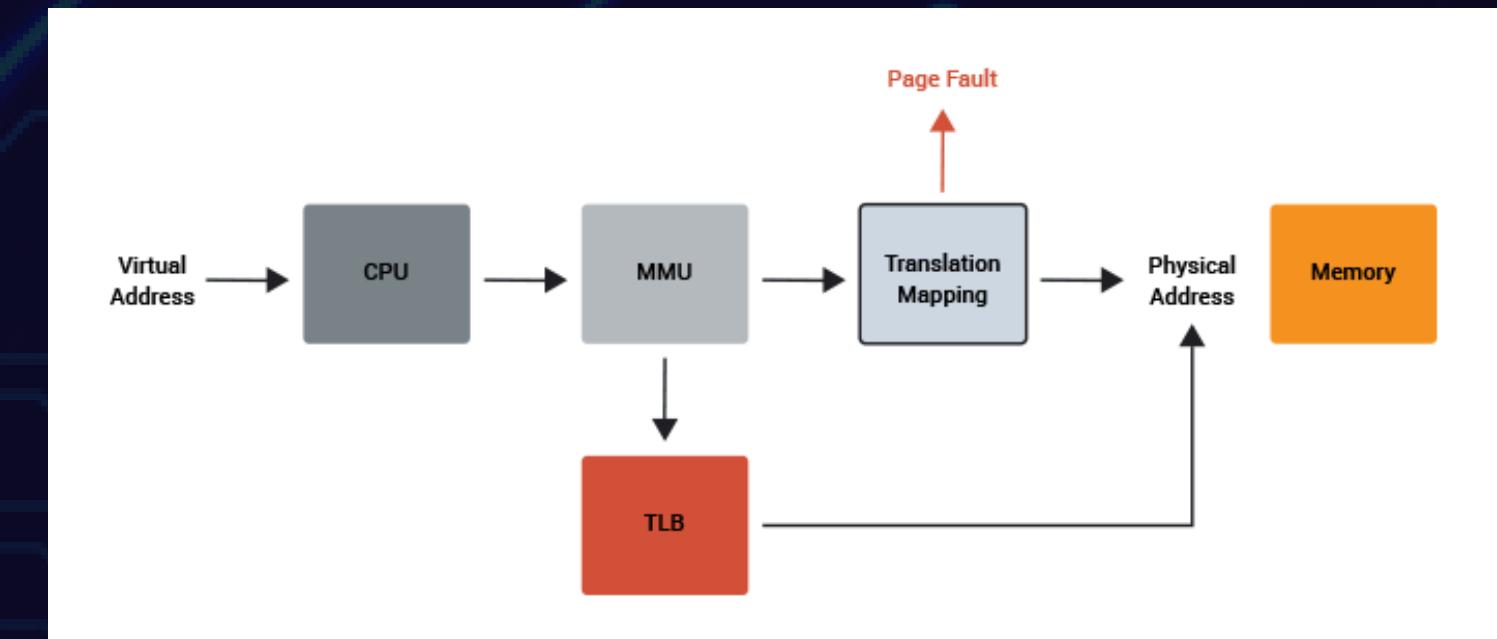
Paging

- Allows memory allocation **to be non-contiguous**
- However, they do **not directly correspond to physical** memory locations.
- Logical address is **translated** into a physical address by the Memory Management Unit (MMU)
- it **provides an abstraction**, allowing programs to think they have access to a large, continuous block of memory when, in reality, physical memory may be fragmented.



Page Fault

- when program accesses a memory page mapped into the virtual address space but not loaded in physical memory
- Since actual physical memory is much smaller than virtual memory, page faults happen.
- The Operating System might have to replace one of the existing pages with the newly needed page.
- similar to *cache miss, if you will..*





IMPLEMENTATION



Paging Algorithms

- Different page replacement algorithms **suggest different ways to decide** which page to replace.
- techniques used in operating systems to manage memory efficiently **when the virtual memory is full**
- Goal for all algorithms is **to reduce the number of page faults.**



Most Common Paging Algorithms

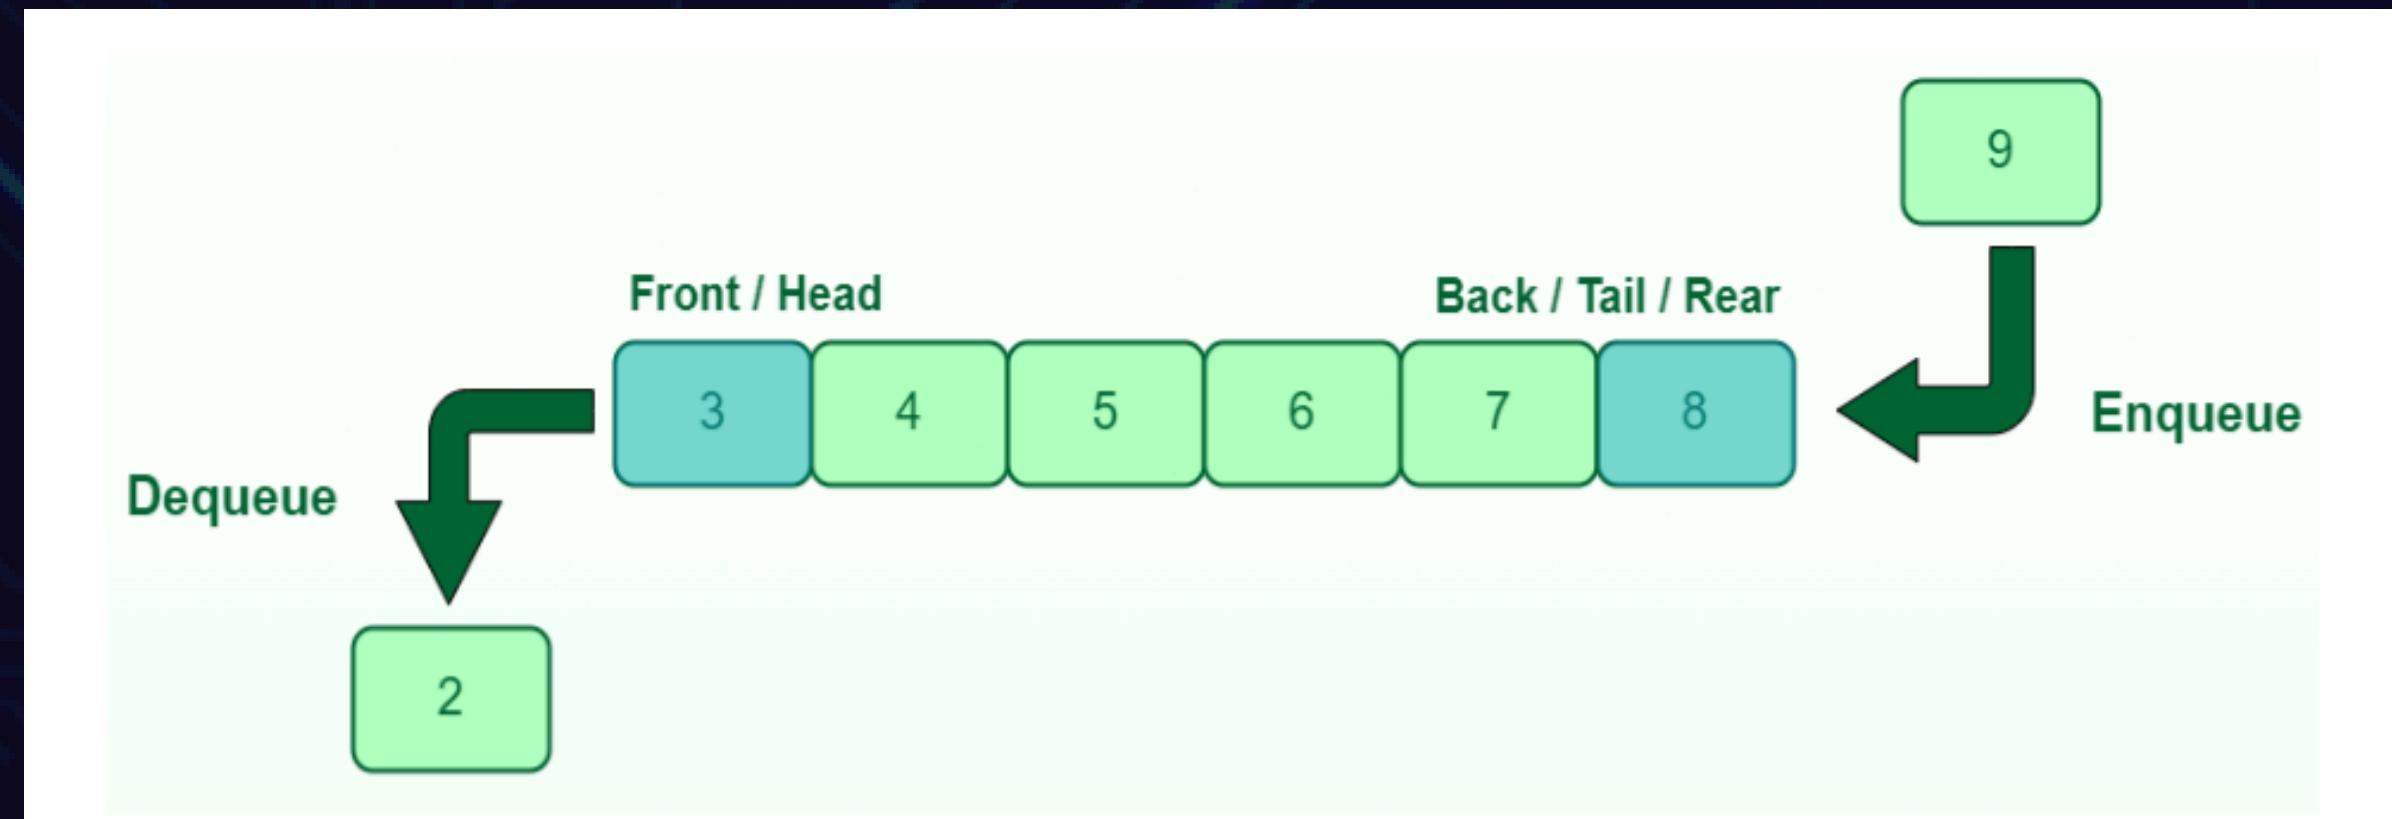
- First In First Out (FIFO).
- Least Recently Used (LRU).
- Optimal Page replacement
- Most Recently Used (MRU)



FIFO

First-In First-Out (FIFO) Algorithm

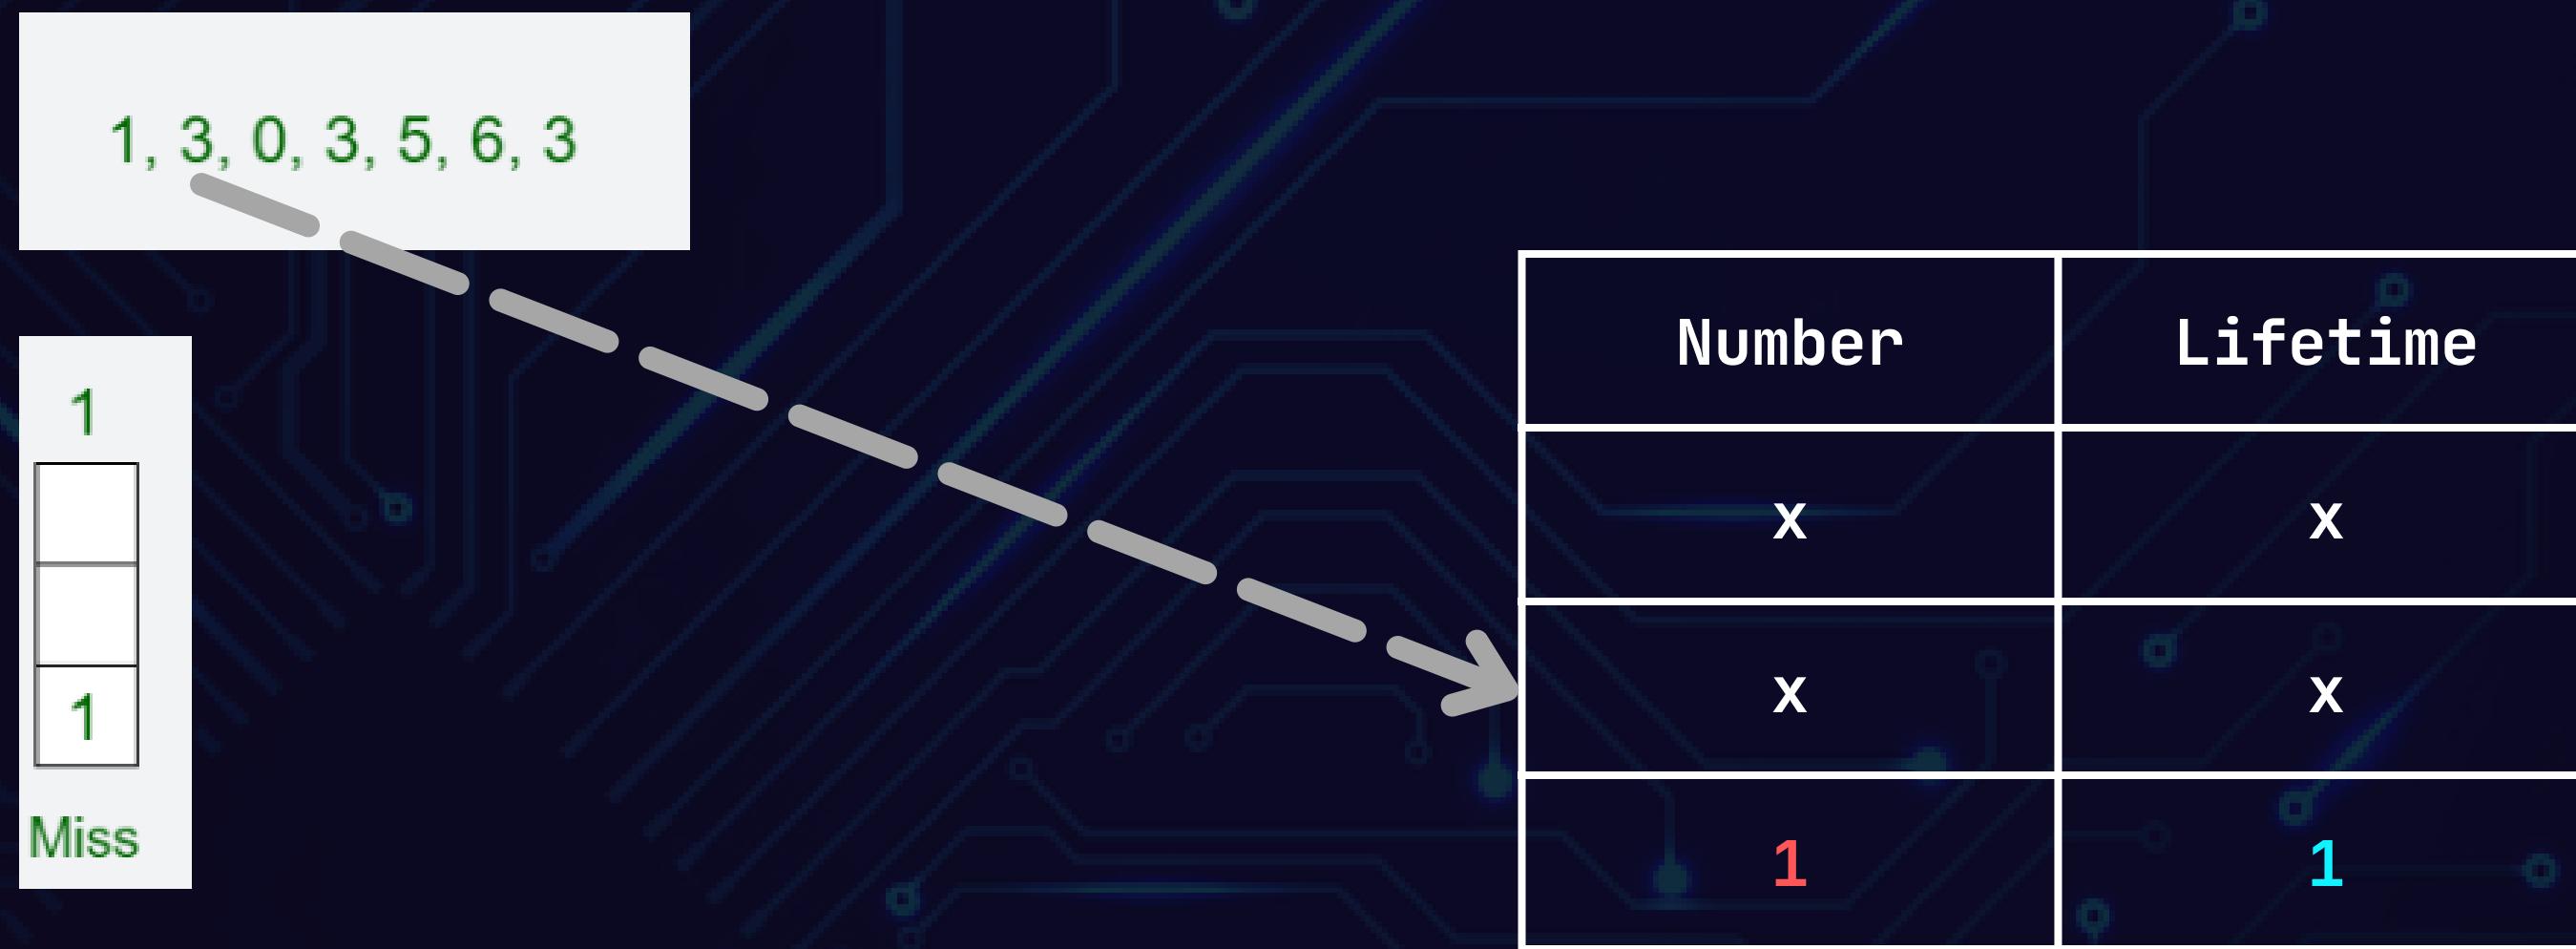
- Queue Data Structure Review



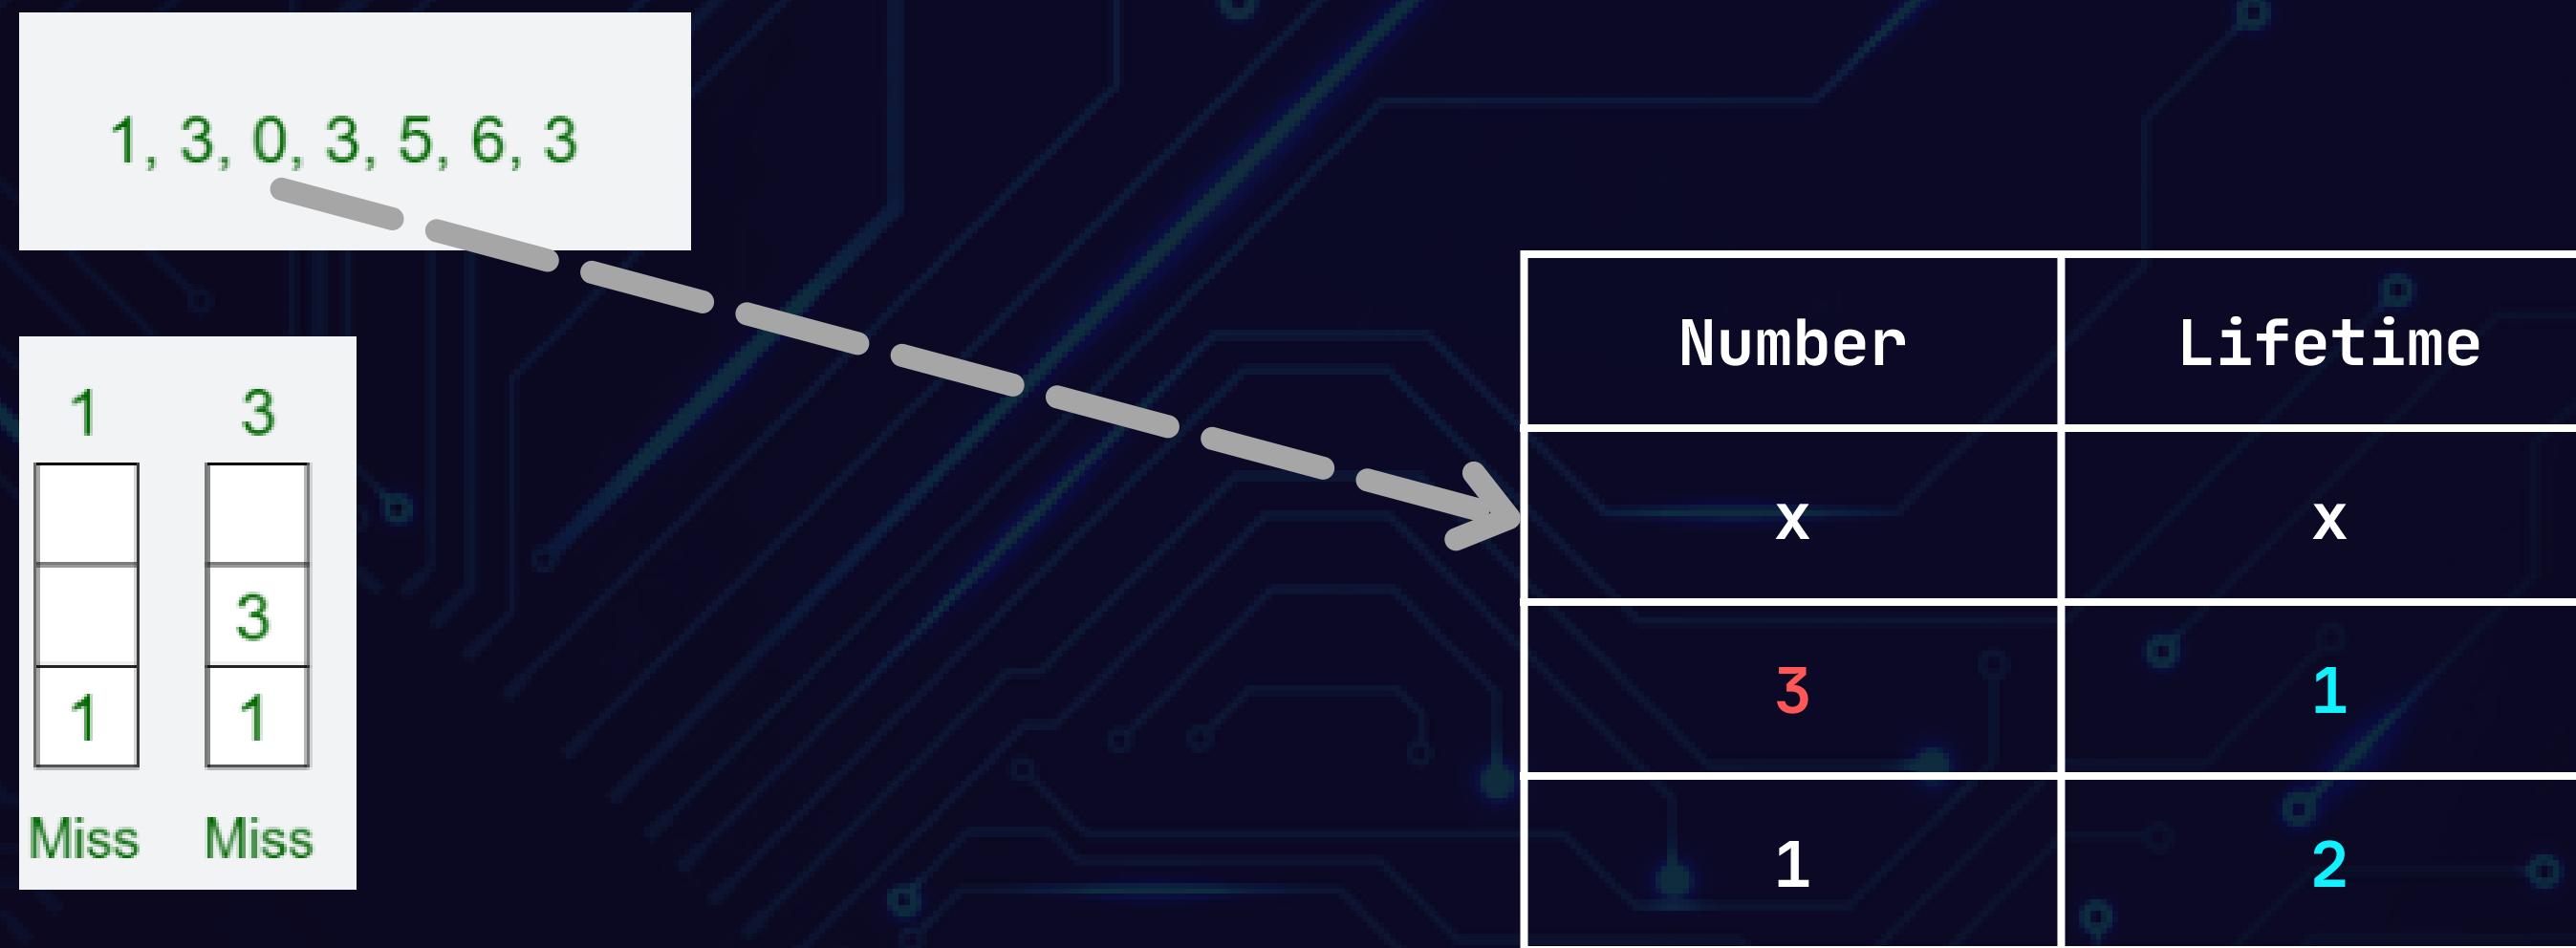


LRU

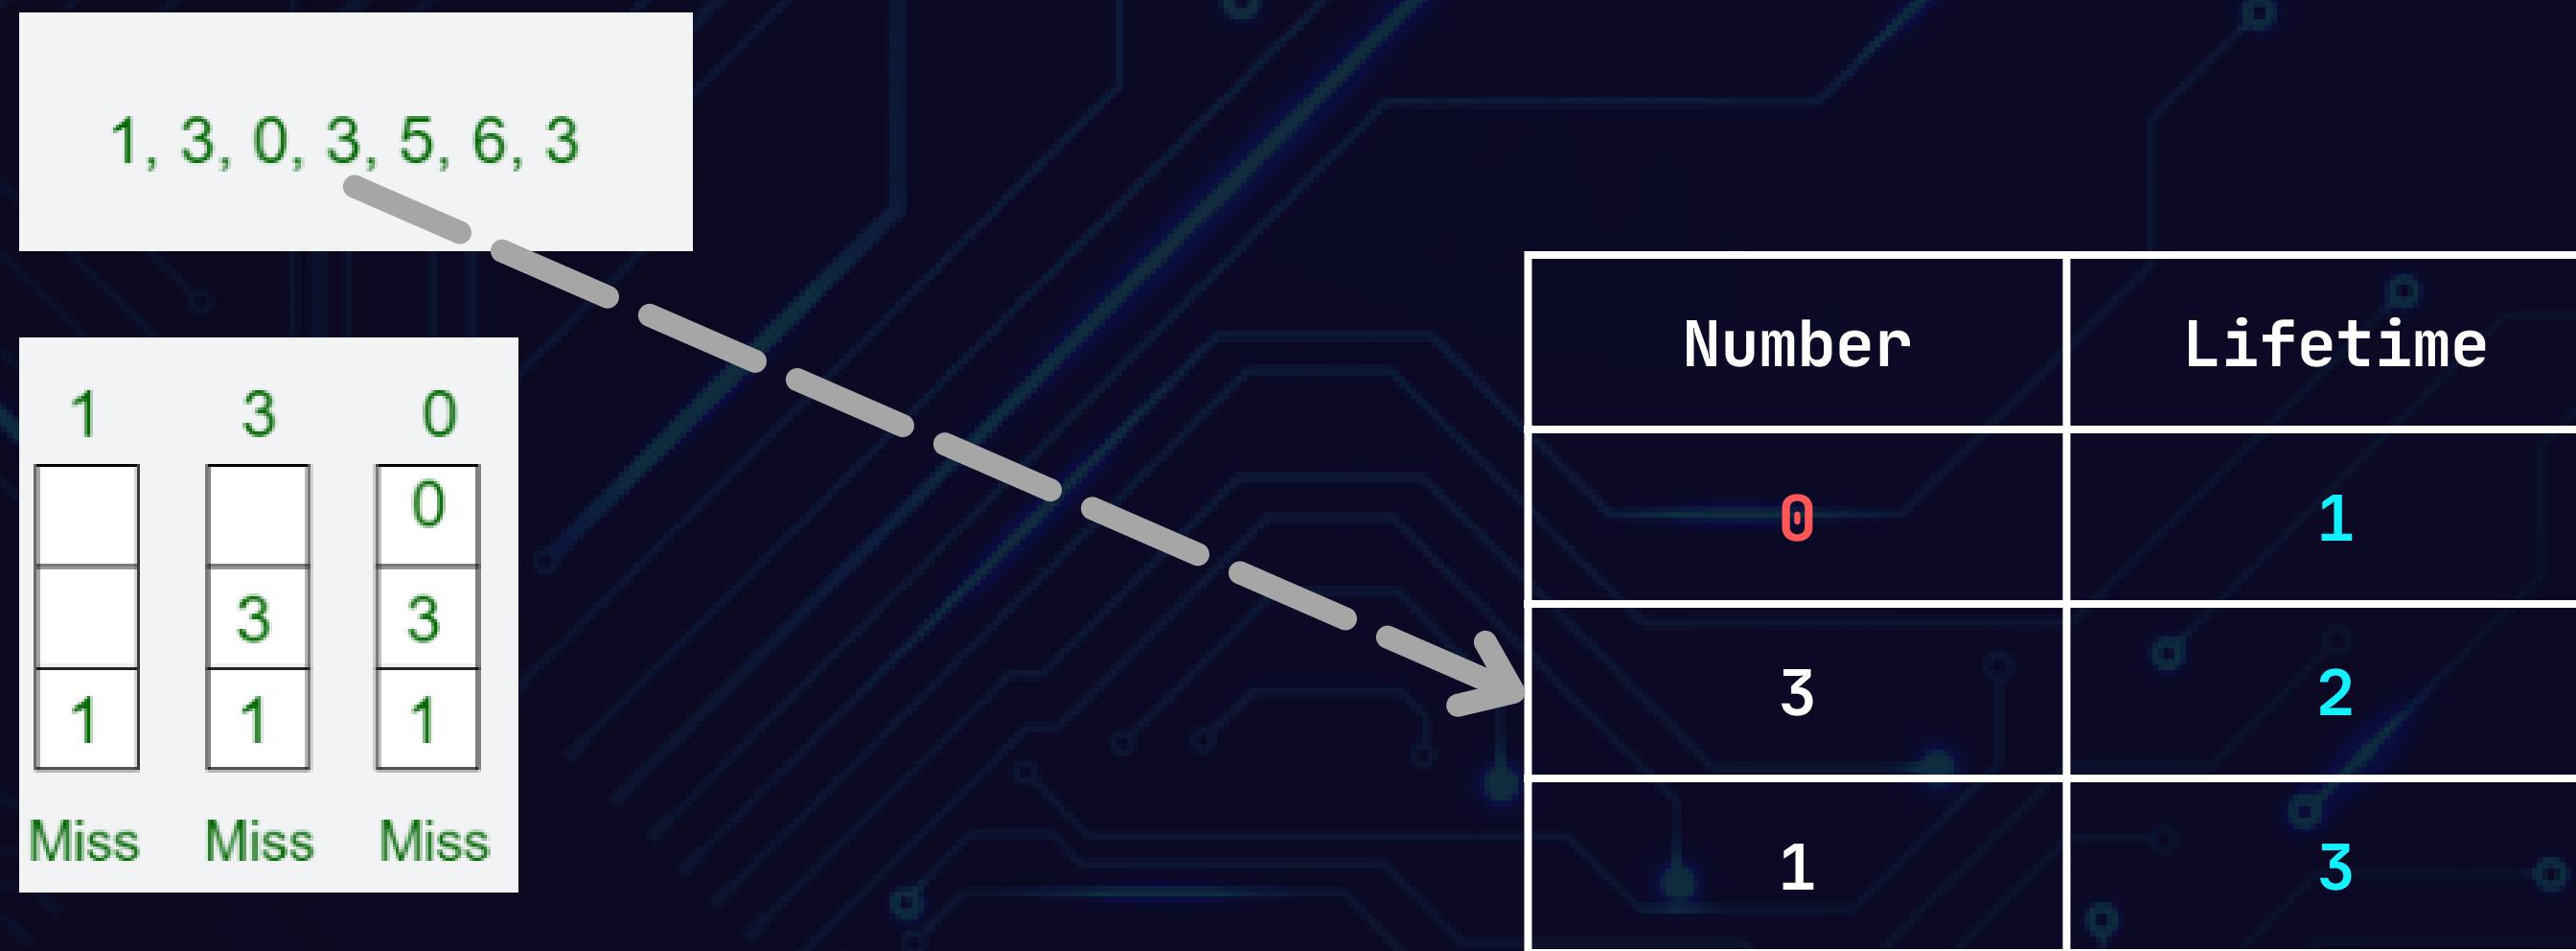
Least Recently Used (LRU) Algorithm



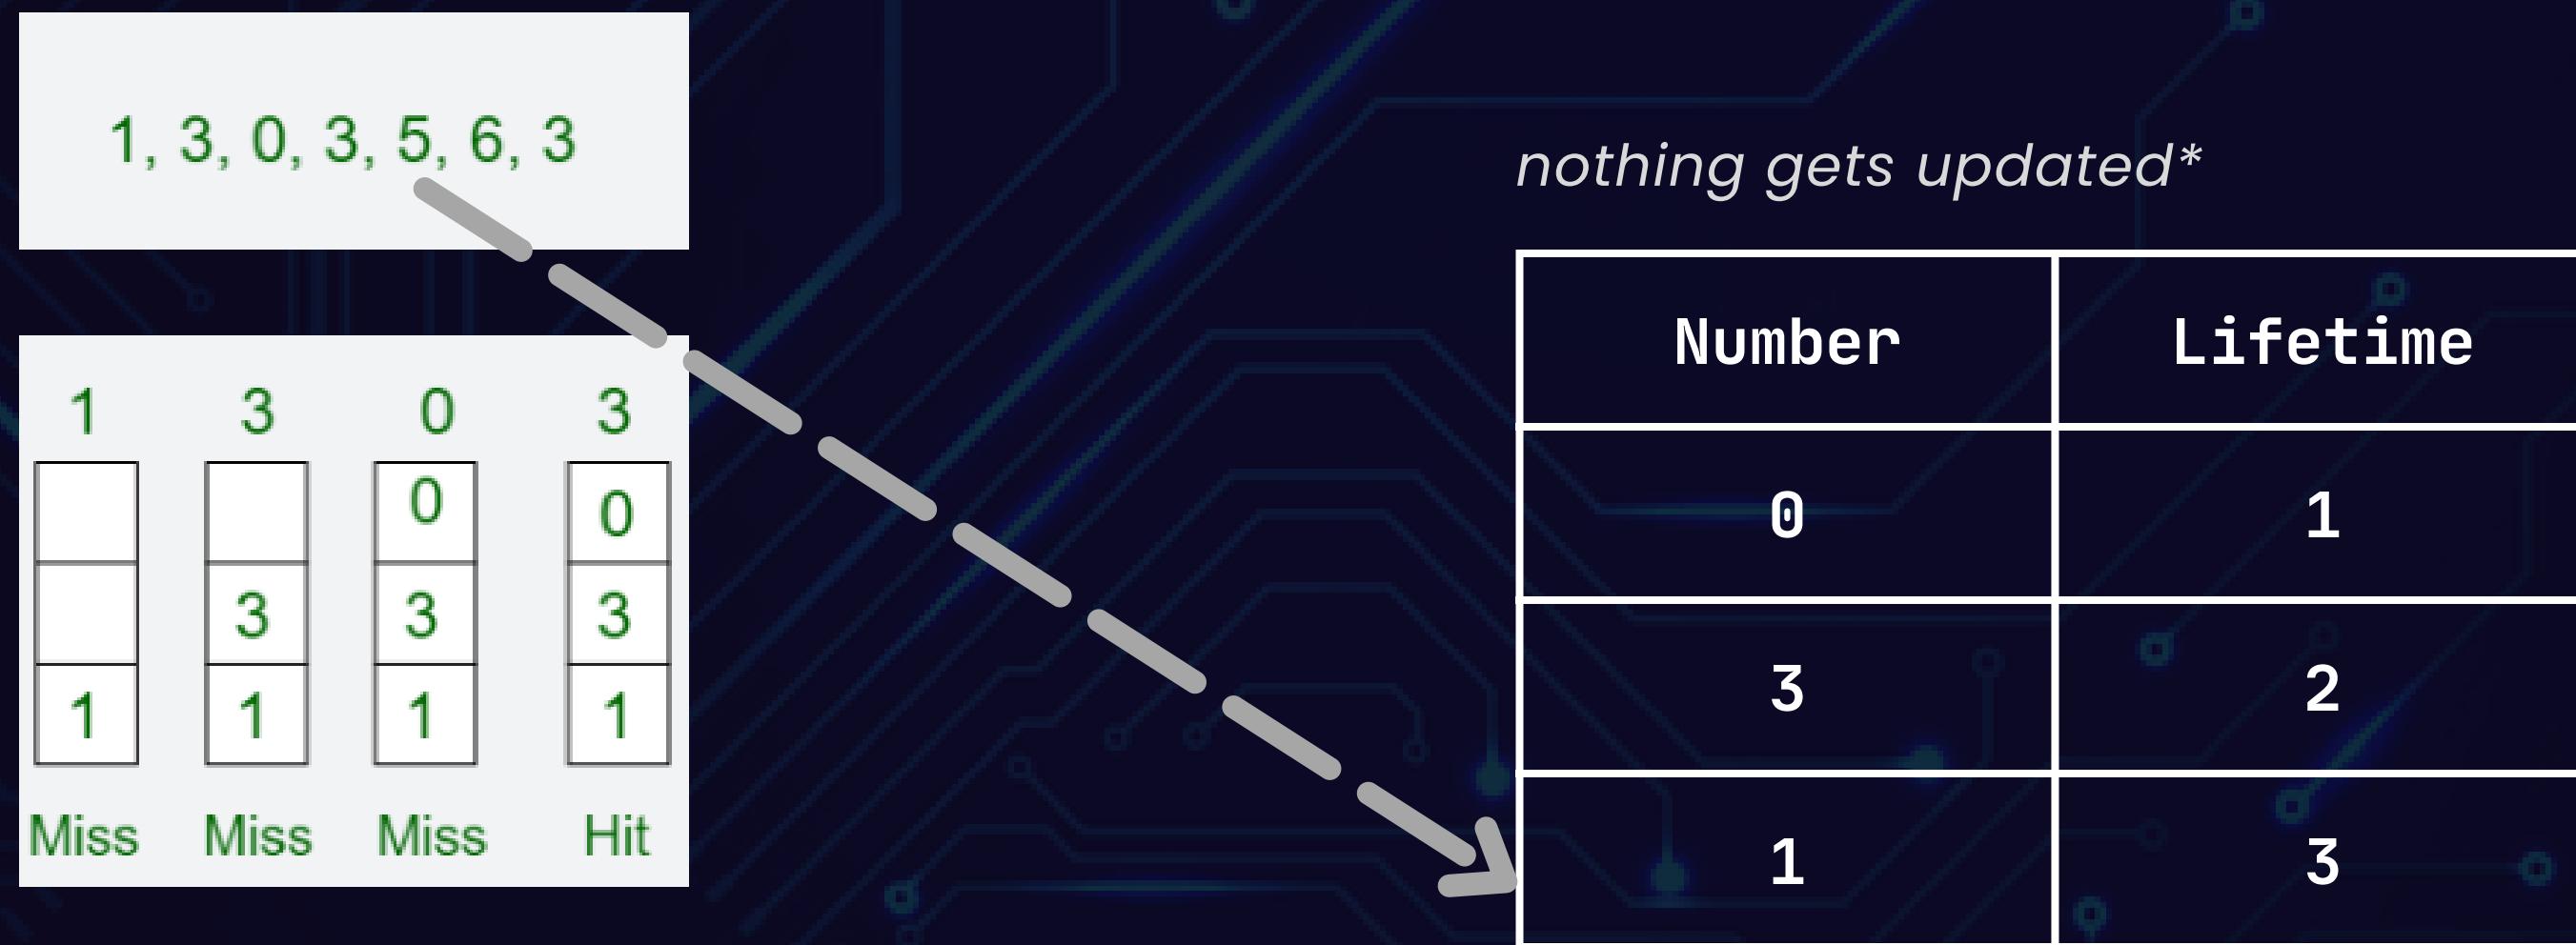
Least Recently Used (LRU) Algorithm



Least Recently Used (LRU) Algorithm



Least Recently Used (LRU) Algorithm



Least Recently Used (LRU) Algorithm

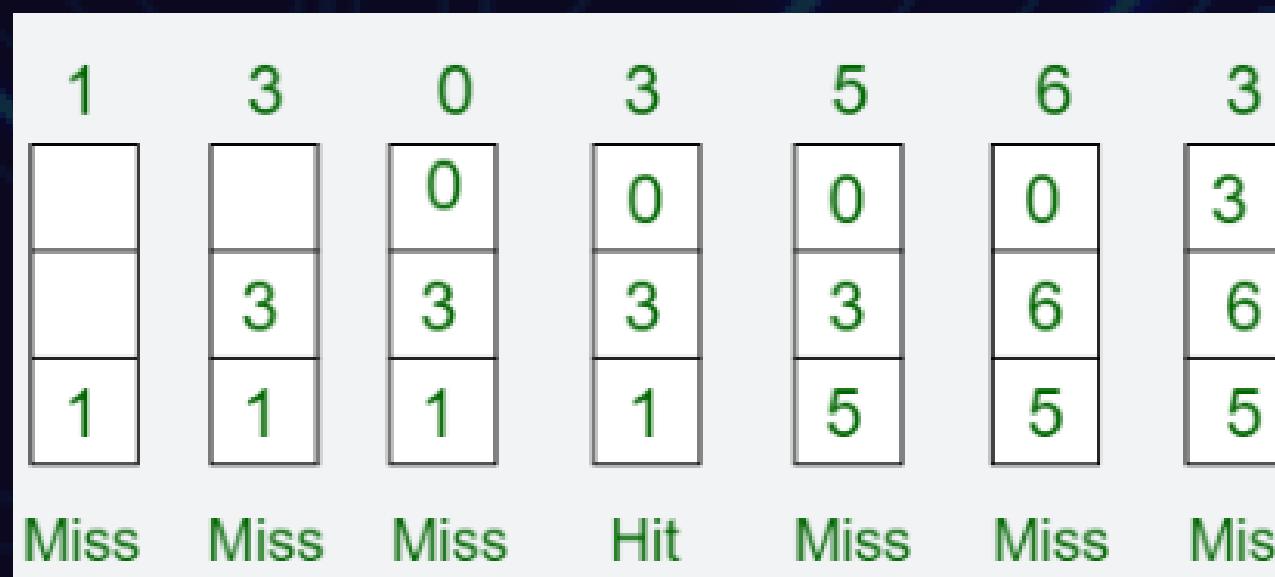


Least Recently Used (LRU) Algorithm



Least Recently Used (LRU) Algorithm

1, 3, 0, 3, 5, 6, 3



done!

Number	Lifetime
3	1
6	2
5	3



Function - First In, First Out (FIFO) Algorithm

```
def fifo_page_replacement(pages: list[int], frame_size: int) -> int:
    memory = [] # List to store pages in memory
    page_faults = 0 # Count of page faults

    for page in pages:
        if page not in memory:
            if len(memory) == frame_size:
                memory.pop(0) # Remove the oldest page (FIFO)
            memory.append(page) # Add the new page
            page_faults += 1
```



Function - Least Recently Used (LRU) Algorithm

```
def lru_page_replacement(pages: list[int], frame_size: int) -> int:
    memory = [] # List to store pages in memory
    page_faults = 0 # To count page faults

    for page in pages:
        if page not in memory:
            if len(memory) == frame_size:
                memory.pop(0) # Remove the least recently used page
            memory.append(page) # Add the new page
            page_faults += 1
        else:
            # Refresh the page by moving it to the end
            memory.remove(page)
            memory.append(page)

    return page_faults
```



Example usage

```
# Example usage
frame_size = 3
reference_string = [1,3,0,3,5,6,3]
page_faults = fifo_page_replacement(reference_string, frame_size)
page_faults = lru_page_replacement(reference_string, frame_size)

print(f"Total page faults: {page_faults}")
```



ANALYSIS



Tests

- Run the FIFO and LRU implementations with various reference strings and frame sizes.
- Record the number of page faults for each scenario.

Tasks

- Test with different reference strings (e.g., random, sequential).
- Compare the performance of FIFO and LRU for the same reference strings.



Testing Procedure

```
# 1. Simple Repeated Pages
pages1 = [1, 1, 1, 1, 1]
frame_sizes1 = [1, 2, 3]

# 2. Sequential Access
pages2 = [1, 2, 3, 4, 5, 6, 7, 8, 9]
frame_sizes2 = [2, 3, 4]

# 3. Alternating Pattern
pages3 = [1, 2, 1, 2, 1, 2, 1, 2]
frame_sizes3 = [1, 2, 3]

# 4. Mixed Pattern
pages4 = [3, 2, 1, 4, 2, 5, 2, 3, 7, 6, 3, 2, 1, 4, 3, 2, 1]
frame_sizes4 = [3, 4, 5]

# 5. Worst-Case for FIFO
pages5 = [1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5]
frame_sizes5 = [3, 4]
```

Testing Results

Test Case	Frame Size	FIFO Faults	LRU Faults
Case 1	1	1	1
Case 1	2	1	1
Case 1	3	1	1
Case 2	2	9	9
Case 2	3	9	9
Case 2	4	9	9
Case 3	1	8	8
Case 3	2	2	2
Case 3	3	2	2
Case 4	3	15	14
Case 4	4	12	10
Case 4	5	11	9
Case 5	3	9	10
Case 5	4	10	8



Testing Function

```
# Function to run test cases and show only page faults
def run_test_cases():
    print(f"{'Test Case':<12}{{'Frame Size':<12}{{'FIFO Faults':<12}{{'LRU Faults':<12}}}")
    print("-" * 48)

    # Test Case 1: Simple Repeated Pages
    for size in frame_sizes1:
        print(f"Case 1    {size:<12}{fifo_page_replacement(pages1, size):<12}{lru_page_replacement(pages1, size):<12}")

    # Test Case 2: Sequential Access
    for size in frame_sizes2:
        print(f"Case 2    {size:<12}{fifo_page_replacement(pages2, size):<12}{lru_page_replacement(pages2, size):<12}")

    # Test Case 3: Alternating Pattern
    for size in frame_sizes3:
        print(f"Case 3    {size:<12}{fifo_page_replacement(pages3, size):<12}{lru_page_replacement(pages3, size):<12}")

    # Test Case 4: Mixed Pattern
    for size in frame_sizes4:
        print(f"Case 4    {size:<12}{fifo_page_replacement(pages4, size):<12}{lru_page_replacement(pages4, size):<12}")

    # Test Case 5: Worst-Case for FIFO
    for size in frame_sizes5:
        print(f"Case 5    {size:<12}{fifo_page_replacement(pages5, size):<12}{lru_page_replacement(pages5, size):<12}")

    # Run all test cases
run_test_cases()
```



DISCUSSION



Performance Comparison

FIFO (First-In-First-Out)	LRU (Least Recently Used)
Advantages	Advantages
Simple to implement	Adaptive to access patterns
Less computationally expensive	Minimizes page faults in many real-world scenarios
Disadvantages	Disadvantages
May evict frequently accessed pages	More complex and requires tracking usage
Not adaptive to access patterns	Slightly higher overhead in memory usage



Performance Comparison

Case 1: Simple Repeated Pages

Frame Size	FIFO Faults	LRU Faults
1	1	1
2	1	1
3	1	1

Analysis:

- Since the page references are all the same ([1, 1, 1, 1, 1]), both FIFO and LRU algorithms incur only one page fault when the first page is accessed. Subsequent accesses do not cause any faults because the page is already in memory. Increasing the frame size does not affect the result in this case because only one unique page is being referenced.

Performance Comparison

Case 2: Sequential Access

Frame Size	FIFO Faults	LRU Faults
2	9	9
3	9	9
4	9	9

Analysis:

- The page references are sequential ([1, 2, 3, 4, 5, 6, 7, 8, 9]). Each new page accessed causes a fault since there's no repetition and the frame cannot hold all unique pages. Both FIFO and LRU algorithms perform identically in this case. No matter the frame size, every new page results in a fault until the end of the sequence, where the frame size is insufficient to hold all unique pages.



Performance Comparison

Case 3: Alternating Pattern

Frame Size	FIFO Faults	LRU Faults
1	8	8
2	2	2
3	2	2

Analysis:

- The pattern `[1, 2, 1, 2, 1, 2, 1, 2]` alternates between two pages.
- With a frame size of 1, each switch between pages causes a fault, resulting in 8 faults for both algorithms.
- With a frame size of 2 or more, the faults drop significantly to just 2, since both pages fit in the frame. Both algorithms perform equally well here as the pattern is simple and predictable.

Performance Comparison

Case 4: Mixed Pattern

Frame Size	FIFO Faults	LRU Faults
3	15	14
4	12	10
5	11	9

Analysis:

- The page references have a mix of sequential and repeated patterns ([3, 2, 1, 4, 2, 5, 2, 3, 7, 6, 3, 2, 1, 4, 3, 2, 1]).
 - **With frame size 3:** Both algorithms have many faults, but LRU performs slightly better, as it can better adapt to the access patterns by replacing the least recently used page.
 - **With frame size 4 or 5:** The number of faults decreases for both algorithms as the frame can hold more pages. LRU still performs better because it efficiently manages recent page accesses.



Performance Comparison

Case 5: Worst-Case for FIFO

Frame Size	FIFO Faults	LRU Faults
3	9	10
4	10	8

Analysis:

- The pattern ([1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5]) is chosen to demonstrate a scenario where FIFO experiences a high number of faults due to the "Belady's anomaly."
 - **With frame size 3:** FIFO has fewer faults than LRU due to the repetitive nature, where LRU mistakenly removes pages too early.
 - **With frame size 4:** LRU performs better, as it better retains recently used pages. FIFO fails to adapt, causing more faults.



Performance Comparison

Summary

- **Case 1:** Both algorithms perform equally well, as there's no need for replacement.
- **Case 2:** Both algorithms perform poorly due to the inability to hold all unique pages.
- **Case 3:** Both perform equally, with a significant reduction in faults as frame size increases.
- **Case 4:** LRU generally performs better with larger frames, as it efficiently manages pages with more diverse patterns.
- **Case 5:** Shows a scenario where LRU can be advantageous, but sometimes FIFO can perform unexpectedly better due to specific patterns causing "Belady's anomaly."

Overall, the choice between FIFO and LRU should depend on the nature of page access patterns and the available frame size. LRU is generally better with varied patterns, while FIFO is simpler and may perform comparably in certain predictable patterns.



DOCUMENTATION



docs.google.com/document/d/1jrJ1Jy2UbkWhgTnwE8jh6dEBfgv5kg_bz...

Memory Management

File Edit View Insert Format Tools Extensions Help

Comments

All comments For you

All types

1. When a new page needs to be loaded into memory and there is no space, the page at the **front of the queue (oldest page)** is replaced first.

Pages >>	6	7	8	9	6	7	1	6	7	8
Frame 3			8	8	8	7	7	7	7	7
Frame 2		7	7	7	6	6	6	6	6	6
Frame 1	6	6	6	9	9	9	1	1	1	1

Reference:

<https://workat.tech/core-cs/tutorial/page-replacement-algorithms>

Opioxh7ym5

- <https://www.youtube.com/watch?v=Z2SL8LmdUSQ>
- <https://www.youtube.com/watch?v=KZcMQjdQf2w>
- Sync Fifo: <https://m.youtube.com/watch?v=FmapuHOfmec>

Comments

1. Joeniño Cainday 6:13 PM Yesterday

2. John

3. John Aaron Sabio 6:59 PM Yesterday

4. Paki check kung ok na ni

5. What does paging mean?

6. Joeniño Cainday 6:15 PM Yesterday

7. John

8. John Aaron Sabio 7:08 PM Yesterday

9. Done na ni

10. John Aaron Sabio 7:08 PM Yesterday

11. Marked as resolved

can you discuss what these results mean

Test Case	Frame Size	FIFO Faults	LRU Faults
Case 1	1	1	1
Case 1	2	1	1
Case 1	3	1	1
Case 2	2	9	9
Case 2	3	9	9
Case 2	4	9	9

Test Case	Frame Size	FIFO Faults	LRU Faults
Case 1	1	1	1
Case 1	2	1	1
Case 1	3	1	1
Case 2	2	9	9
Case 2	3	9	9
Case 2	4	9	9

Message ChatGPT

Roadmap | Memory M | Memo | Facebook | Page Repla | (51) The Cl | Operating | Google paging me | Google caching to | ChatGPT | (51) LRU (| ChatGPT | +



Explorer

main.py > ...

OS-REPORT

main.py

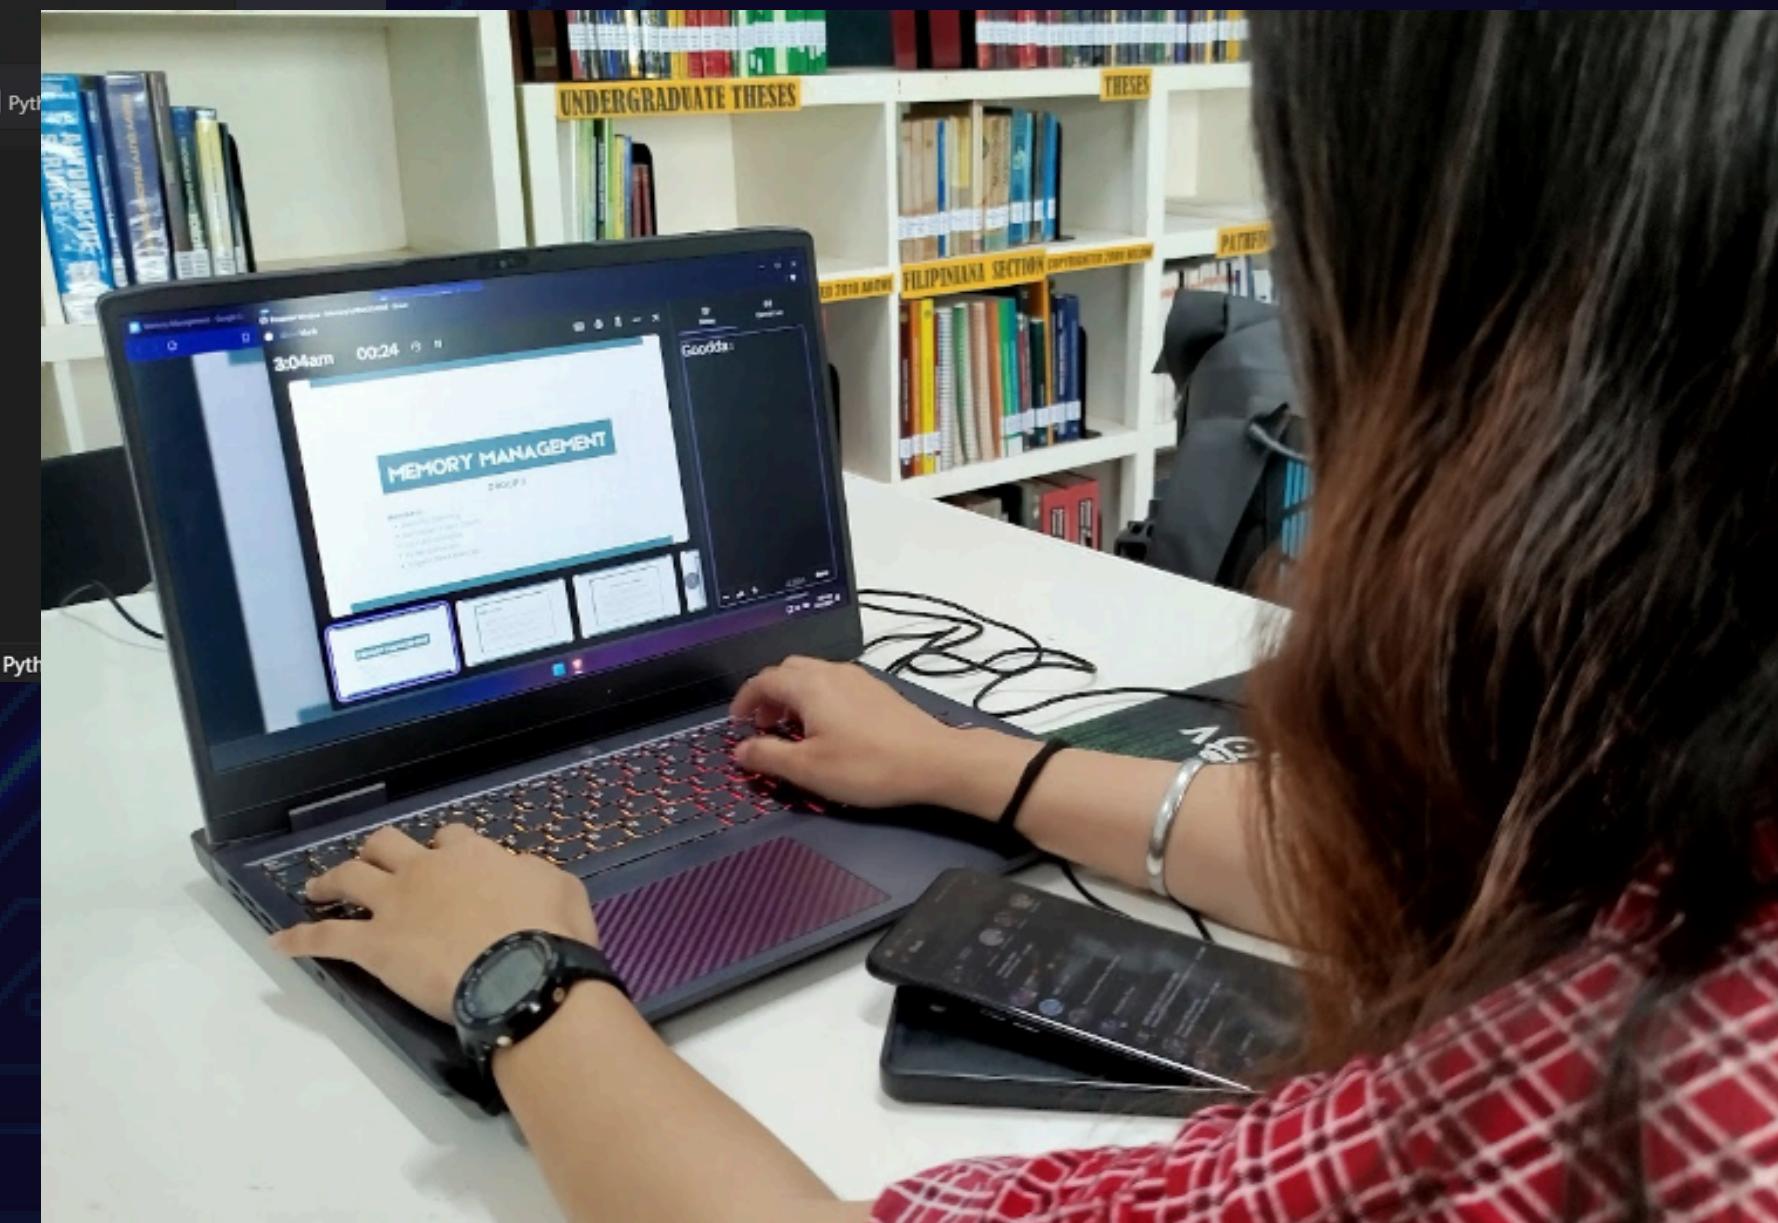
```
1 def fifo_page_replacement(pages: list[int], frame_size: int) -> int:
2     memory = [] # List to store pages in memory
3     page_faults = 0 # Count of page faults
4
5     for page in pages:
6         if page not in memory:
7             if len(memory) == frame_size:
8                 memory.pop(0) # Remove the oldest page (FIFO)
9             memory.append(page) # Add the new page
10            page_faults += 1
11
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS COMMENTS

Case	2	3	9	9
Case 2	4	9	9	9
Case 3	1	8	8	8
Case 2	4	9	9	9
Case 3	1	8	8	8
Case 3	1	8	8	8
Case 3	2	2	2	2
Case 3	3	2	2	2
Case 4	3	15	14	
Case 3	3	2	2	
Case 4	3	15	14	
Case 4	3	15	14	
Case 4	4	12	10	
Case 4	5	11	9	
Case 5	3	9	10	

Ln 80, Col 1 Spaces: 4 UTF-8 CRLF {} Python

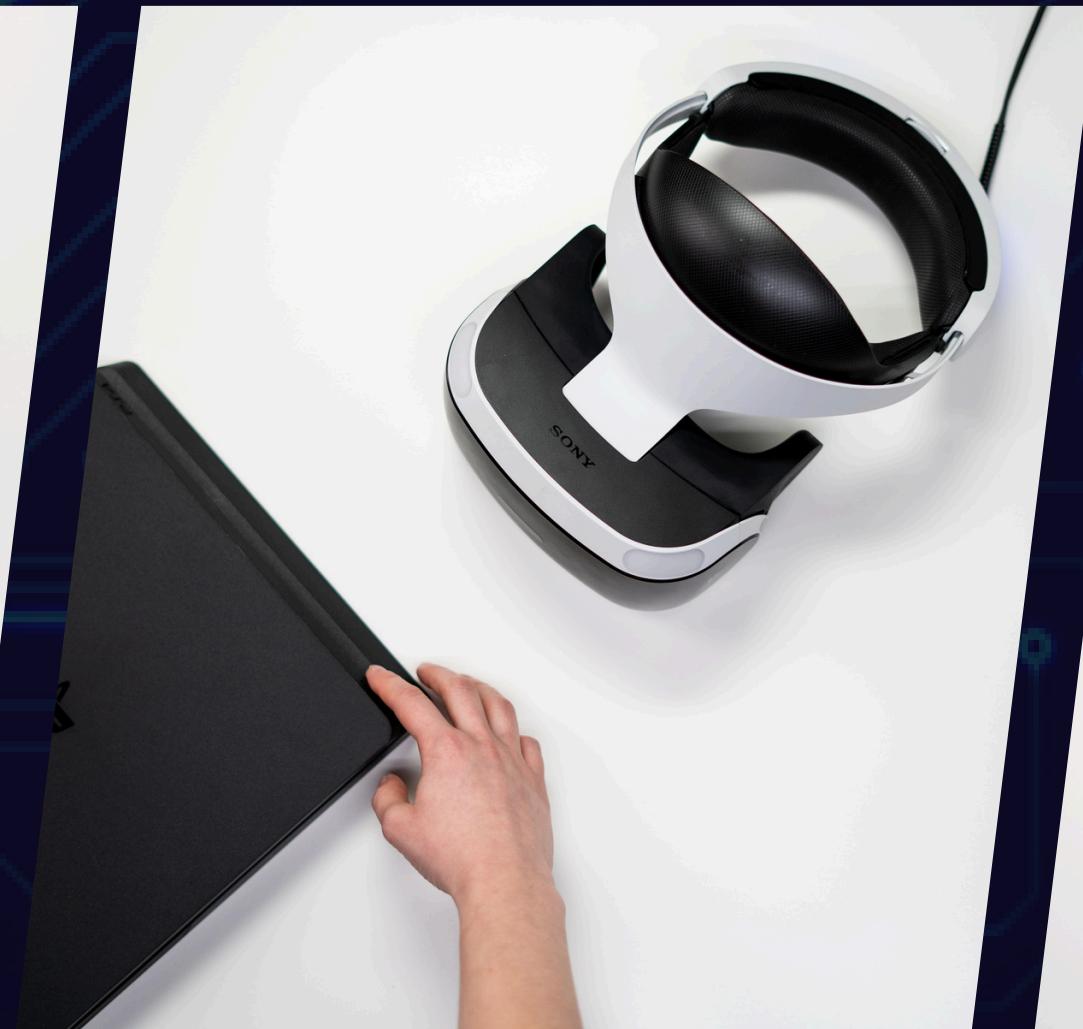
Live Share





REFERENCES

- tutorialspoint.com/operating_system/os_memory_management.htm
- wikipedia.org/wiki/Memory_management
- [wikipedia.org/wiki/Memory_management_\(operating_systems\)](http://wikipedia.org/wiki/Memory_management_(operating_systems))
- courses.cs.washington.edu/courses/cse451/04wi/projects/mmlinux.html
- techtarget.com/searchstorage/definition/address-space
- geeksforgeeks.org/page-replacement-algorithms-in-operating-systems/
- geeksforgeeks.org/lazy-swapping-in-operating-system/





MEET OUR TEAM



Rozel Galceran



Angela Mendez



John Aaron Sabio



Nomeben Clarin



Joeninyo Cainday



THANK YOU FOR LISTENING